

# Whirlwind tour of Scala Native

Denys Shabalin, LAMP/EPFL

October 13, 2016

# Scala?

- First appeared on January 20, 2004.
- Has grown to become a prominent JVM language.

# Scala?

- *“At the root, the language’s scalability is the result of a careful integration of object-oriented and functional language concepts.”*

<http://www.scala-lang.org/what-is-scala.html>

# Scala?

- Key insight: *functional programming is cool and we should have more of that in object-oriented languages.*
- Support for first-class functions, higher kinded types, pattern matching, local type inference, tuples, immutable collections etc etc.
- While preserving tight interoperability with Java.

# JVM-only?

- There have been quite a few mixed-results experiments that tried to make it work off the JVM over the years:
  - Scala .NET (2011, Miguel Garcia)
  - Scala GWT (2012, Grzegorz Kossakowski)
  - Scala LLVM (2013, Geoff Reedy)

# Scala.js

- First successful alternative platform: Scala.js (2013, Sébastien Doeraene)
- In-house incremental whole-program optimising AOT compiler to JavaScript.

# Scala.js

- Introduced its own tree-based IR as a byte-code format.
- Advanced whole-program tree shaking and plethora of Scala-specific optimisations (e.g. aggressive closure-biased inliner.)
- Support for subset of JDK libraries.
- Seamless interoperability with JavaScript.

# Scala Native

- Whole-program optimising AOT compiler to native code built on top of LLVM.
- Easy interoperability with C ABI-compatible code.




# Scala Native

The screenshot shows the GitHub repository page for `scala-native / scala-native`. The repository has 147 watchers, 2,082 stars, and 96 forks. It contains 607 commits, 6 branches, 0 releases, and 30 contributors. The license is BSD-3-Clause. The repository description is "Your favourite language gets closer to bare metal. <http://scala-native.org> — Edit". The repository is currently on the `master` branch. A recent commit by `densh` merged pull request #322 from `Yoitsumi/topic/fix-varargs` is shown. The commit message is "Merge pull request #322 from Yoitsumi/topic/fix-varargs". The commit hash is `c857f6a` and it was made 6 hours ago. The commit details table lists the following changes:

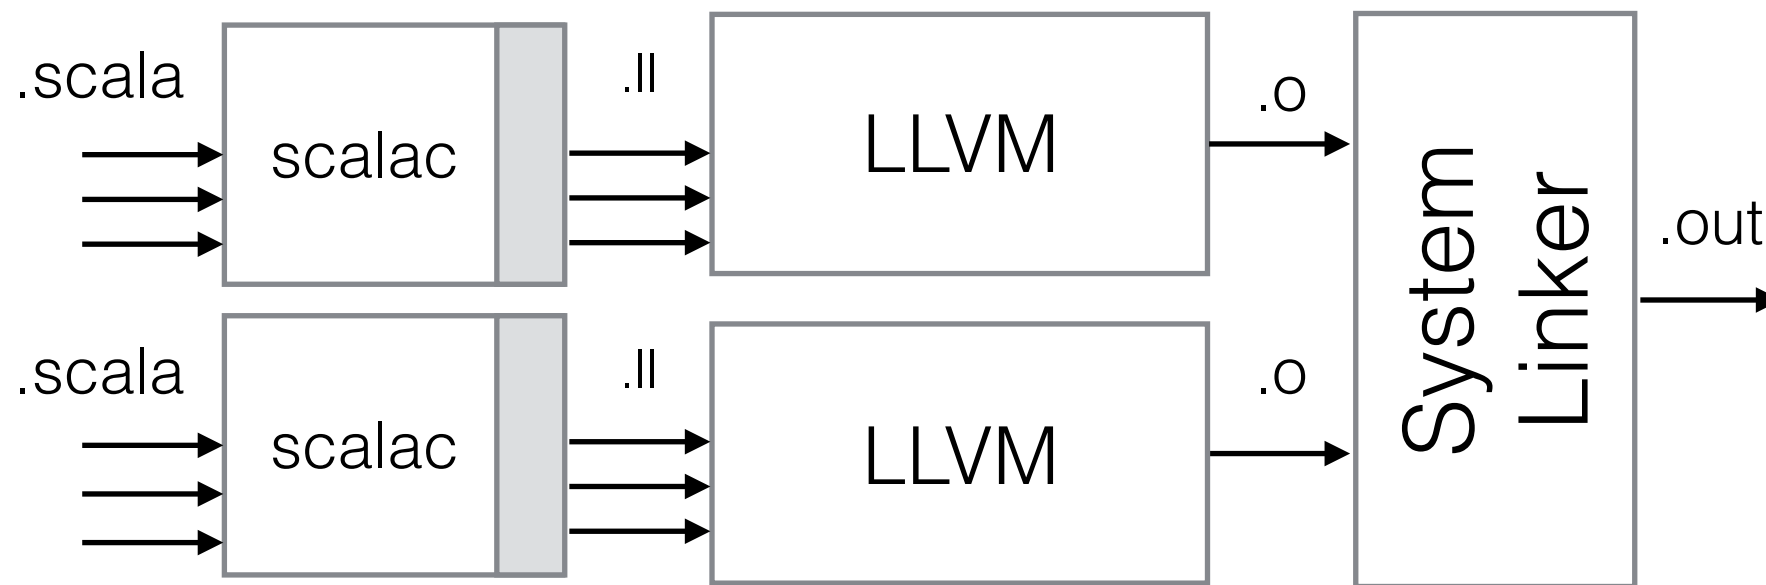
File	Commit Message	Time Ago
<code>demo</code>	Update demo progress indicator on same line	a month ago
<code>docker</code>	Fix #102	4 months ago
<code>docs</code>	Update sbt.rst: add project/build.properties	2 days ago
<code>javalib/src/main/scala</code>	Add toString impl to java.lang.Enum	9 days ago
<code>nativelib/src/main</code>	Modularise the runtime code	6 days ago
<code>nir/src/main/scala/scala/scalanative/...</code>	Introduce function argument attributes	6 days ago
<code>nscplugin/src/main</code>	Introduce function argument attributes	6 days ago
<code>project</code>	Setup scripted tests	3 months ago
<code>sandbox</code>	Fix c-style varargs accepting only single parameter	a day ago
<code>sbtplugin/src/main/scala/scala/scala...</code>	Add support for C files in runtime	6 days ago
<code>scalalib/overrides-2.11/scala</code>	implement array.clone() and basic array.copy() (3)	4 months ago

How does it work?

# Clang



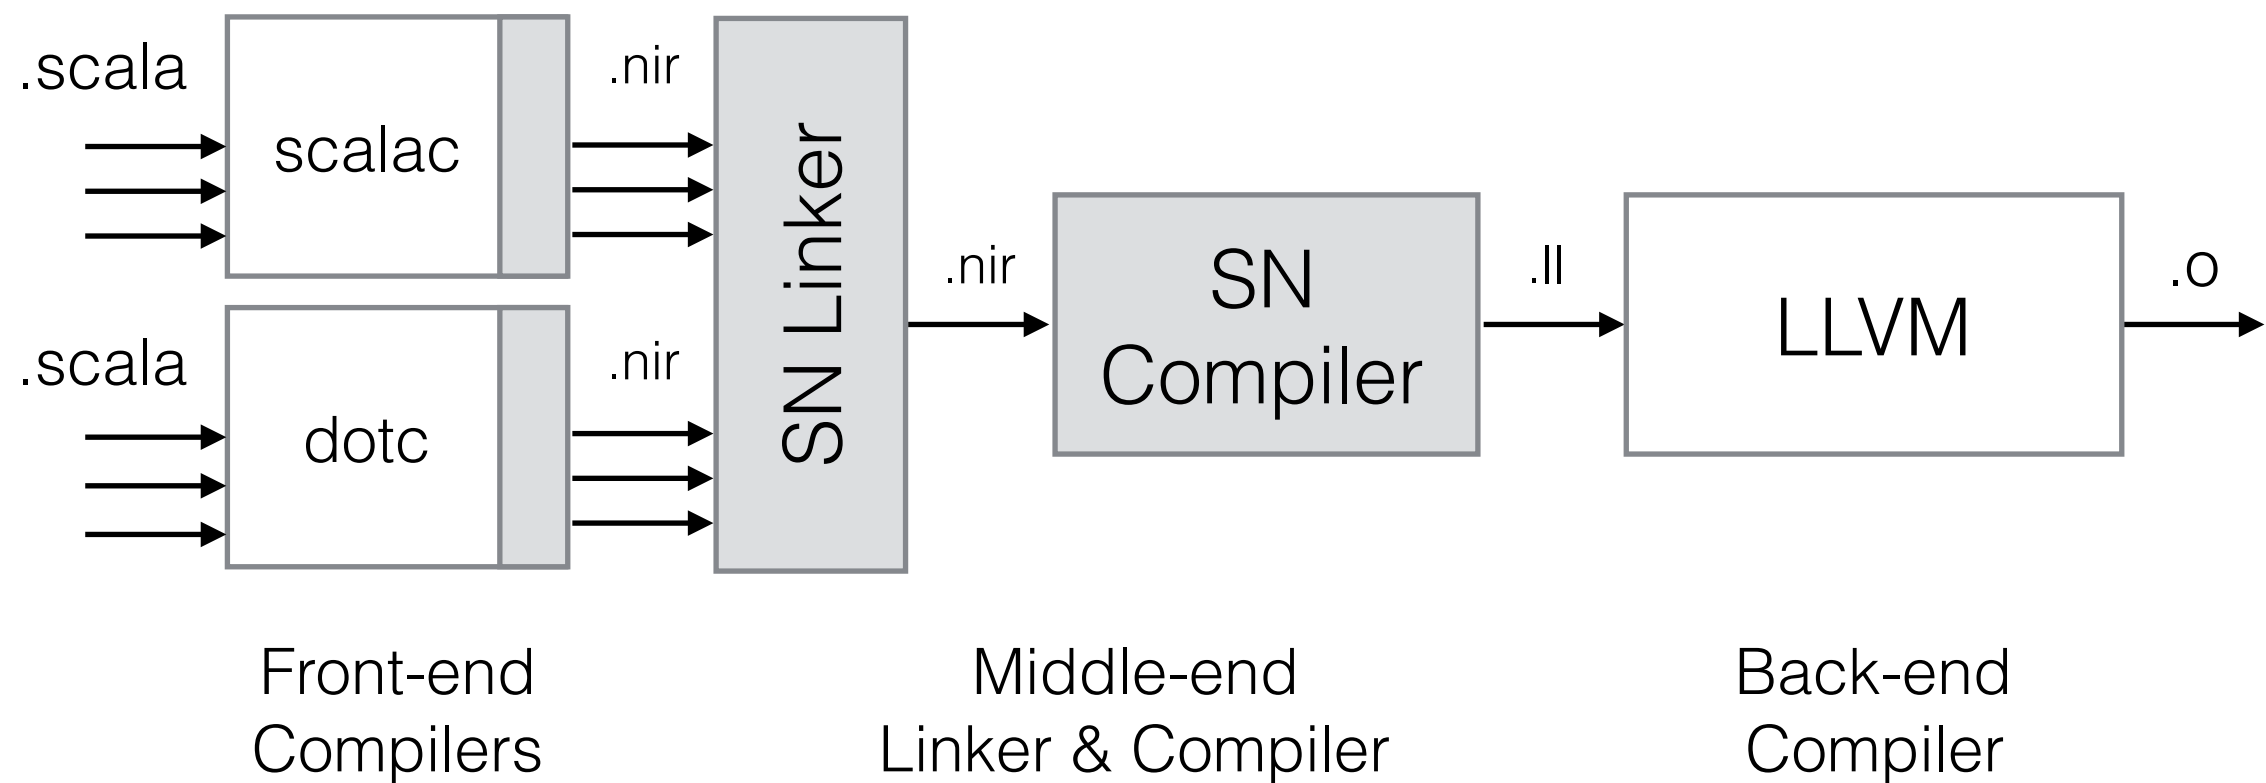
# Scala LLVM



# Problems

- How do you portably distribute compiled code?
  - LLVM IR?
  - Native Object Files?
- How do you optimise features that LLVM doesn't really know about across module boundaries?

# Scala Native



# Front-ends (2x)

- One stable and battle-tested and one experimental and exciting.
- Parses, type-checks, infers types.
- Translates away all the functional features to corresponding object-oriented equivalents.
- Mostly the same codebase for JVM, JS and Native back-ends.

# Middle-end


- Link and tree shake the whole program.
- Compile all the higher-level features away given the static knowledge of the remaining code.
- Single intermediate representation: NIR



# NIR

- Typed high-level object-oriented SSA representation.
- Mixes high-level and low-level primitives.
- We took quite a bit of ideas from SIL.

# NIR



# NIR

How to draw an owl

1.



1. Draw some circles

2.



2. Draw the rest of the owl

# NIR

```
object Test {  
  def main(args: Array[String]): Unit =  
    println("Hello, world!")  
}
```

# NIR

```
module @Test$ : @java.lang.Object

def @Test$::init : (module @Test$) => unit {
  %src.1(%src.0: module @Test$):
    %src.2 = call[...] @java.lang.Object::init(%src.0)
    ret unit
}

def @Test$::main_class.ssr.0bjectArray_unit : ... {
  %src.2(%src.0: module @Test$,
    %src.1: class @scala.scalanative.runtime.ObjectArray):
  %src.3 = module @scala.Predef$
  %src.4 = method[...] %src.3,
    @scala.Predef$::println_class.java.lang.Object_unit
  %src.5 = call[...] %src.4(%src.3, "Hello, world!")
  ret %src.5
}
```

# Middle-end

- Most of NIR trivially maps to LLVM.
- Apart from a few problematic areas.

# Pain points

- First-class lazy modules
- Virtual method dispatch
- Value types & boxing
- Compilation time
- Garbage Collection

# First-class lazy modules

- No top-level members in Scala
- Modules can extend classes and traits
- Initialisation happens on first access



# First-class lazy modules

```
object Test {  
  def main(args: Array[String]): Unit =  
    println("Hello, world!")  
}
```

# First-class modules

```
@"value.scala.Predef$" = global i8* zeroinitializer
```

```
// check if @value.scala.Predef == null
```

```
// if not return old value
```

```
// otherwise allocate & initialise new instance
```

```
define i8* @"load.scala.Predef$"() { ... }
```

```
define void @"Test$::main_class.ssnr.ObjectArray_unit"(i8* %src.0, i8* %src.1) {  
src.2:
```

```
    %src.3 = call i8* () @"load.scala.Predef$"()
```

```
    call void (i8*, i8*) @"scala.Predef$::println_class.java.lang.Object_unit"(  
        i8* %src.3, i8* bitcast ({ i8*, i32, i32, i32, i8* }* @"__const.1" to i8*))
```

```
    ret void
```

```
}
```

# First-class modules

- Short-term solutions:
  - Don't generate accessors for stateless objects with side-effect-free constructors
  - Use GVN over NIR to eliminate duplicate method loads

Under development by Hugo Kapp.

# Virtual methods

- Analyse class hierarchy and detect effectively final methods
- Use vtables for class virtual method dispatch
- Use statically generated dispatch table for interface virtual method dispatch

# Virtual methods

- On-going experiments:
  - Off-line type profiling
  - Online type profiling and inline caching in AOT setting with LLVM's patch-points.

Under research by Martin Duhem.

# Value types & boxing

```
object Test {  
    def main(args: Array[String]) = {  
        val elems = (1 to 10000).toArray  
        val transformed = elems.map(_ * 2)  
        transformed.foreach(println)  
    }  
}
```

- Allocates two integer boxes on every iteration in map
- Allocates an integer box on every iteration in foreach

# Value types & boxing

- Simple solution:
  - Pack values into pointers (aka SMI).
  - 62 bits of space should be enough for everyone.
  - Boxing/unboxing is going to shuffle bits around without actually allocating anything.

# Compilation time

- Solution:
  - Incremental & parallel middle-end.
  - Incremental back-end with ThinLTO.

“Parallel Incremental Whole-Program Optimizations for Scala.js”

by Sébastien Doeraene to appear in OOPSLA’16

<http://lampwww.epfl.ch/~doeraene/publications/oopsla16-incremental-optimizer.pdf>



# Garbage collection

- SN currently uses Boehm GC.
- No open-source high-quality LLVM-compatible GCs available to date.
- Azul seems is upstreaming improvements in this area.

# Garbage collection

- Likely solution: write our own.
- Pitfalls:
  - There is no ultimate solution to root extraction (statepoint vs gcroot vs conservative)
  - LLVM optimising things in GC-unsafe way (see strong GC references discussion for details)

# Questions?