

Fast startup and low latency: pick two

Denys Shabalin and Lukas Kellenberger

EPFL

Scala Native

- Announced a year ago with a first prototype at ScalaDays '16 in New York.
- An ahead-of-time compiler for Scala build on top of the LLVM compiler infrastructure.
- Developed by EPFL and Scala Center.

Scala Native

- As of today:
 - 55 contributors
 - 383 pull requests closed
 - 246 issues fixed

Scala Native

Thanks to everyone who contributed!

Guillaume Massé, Martin Duhem, Hugo Kapp, Hiroyoshi Takahashi, Jonas Fonseca, Lukas Kellenberger, Francois Bertrand, AndreaTP, Eric K Richardson, Marius B. Kotsbak, Kota Mizushima, Łukasz Indykiewicz, Timothy Klim, Paweł Batko, Shunsuke Otani, Nick Pavlov, Cedric Viaccoz, Andrzej Sołtysik, Ankit Soni, Hanns Holger Rutz, Kamil Tomala, Philipp Dörfler, Simon Ochsenreither, Zack Powers, Musabilal, Hubert Plociniczak, Mike Samsonov, Greg Dorrell, Pablo Guerrero, Florian Duraffourg, Alex Dupre, Ragavendar Ramamurthi, Richard Whaling, Roman Zoller, Ruben Berenguel Montoro, Saleem Ansari, Sam Halliday, Felix Mulder, Ragnr, Stefan Ollinger, The Gitter Badger, Tim Nieradzick, Brad Rathke, Viacheslav Blinov, Vincent Munier, Remi, Alexey Kutepov, Adam Voss, Gregor-i, Gute-ist-tot, Kenji Yoshida, Joseph Price, Jentsch, Ignat Loskutov, Martin Mauch

Road towards 0.1

(March 14, 2017)

Goals for 0.1

- All Scala language features are supported
- Sbt integration is sufficient to build and publish existing cross-platform projects
- Enough core libraries to cover for basic standard library usage

Improving the standard library story in 0.2

(April 26, 2017)

Goals for 0.2

- Support for file i/o from `java.io.*`
- Support for regex from `java.util.regex.*`
- Event-loop-based Futures

Laying down the foundation for *better* garbage collection in 0.3


(June 6, 2017)

Goal for 0.3

Better GC #128



ebrucez opened this issue on May 23, 2016 · 21 comments



ebrucez commented on May 23, 2016



Adding this as a backlog idea.

[sngen](#) is Mono's newish garbage collector, introduced to replace the old Boehm. It is a generational gc, and since it works for Mono, I would think it could fit the bill for Scala Native as well.

Here is the [source code](#). It appears to under the MIT license.



5

`nativeGC := "boehm"`

Boehm GC

- Conservative garbage collector
- Originally designed for C/C++ environment
- Good starting point for a new language implementation due to its simple GC interface

Conservative GC

- Conservative roots:

GC doesn't *precisely* know which values on the stack are heap references, but object layout is known.


- Fully conservative:

GC doesn't *precisely* know which values on the stack or heap are references.

Boehm GC

- How fast/slow are we exactly?
- github.com/smarr/are-we-fast-yet

Boehm GC



`nativeGC := "none"`

Running without GC

- The simplest form of garbage collection: allocate and never free allocated memory
- Practical for short-lived command-line tools
- Sometimes used in applications with insane requirements for application predictability

From: k...@rational.com (Kent Mitchell)
Subject: Re: Does memory leak?
Date: 1995/03/31
newsgroups: comp.lang.ada

This sparked an interesting memory for me. **I was once working with a customer who was producing on-board software for a missile.** In my analysis of the code, I pointed out that they had a number of problems with storage leaks. Imagine my surprise when the customer's chief software engineer said "Of course it leaks". He went on to point out that they had calculated the amount of memory the application would leak in the total possible flight time for the missile and then doubled that number. **They added this much additional memory to the hardware to "support" the leaks. Since the missile will explode when it hits its target or at the end of its flight, the ultimate in garbage collection is performed without programmer intervention.**

--

Kent Mitchell	One possible reason that things aren't
Technical Consultant	going according to plan is
Rational Software Corporation	that there never <i>was</i> a plan!

<https://groups.google.com/forum/message/raw?msg=comp.lang.ada/E9bNCvDQ12k/1tezW24ZxdAJ>

Running without GC



[OpenJDK FAQ](#)
[Installing](#)
[Contributing](#)
[Sponsoring](#)
[Developers' Guide](#)

[Mailing lists](#)
[IRC · Wiki](#)

[Bylaws · Census](#)
[Legal](#)

JEP Process

Source code

[Mercurial](#)
[Bundles \(6\)](#)

Groups

[\(overview\)](#)
[2D Graphics](#)
[Adoption](#)
[AWT](#)
[Build](#)
[Compiler](#)
[Conformance](#)
[Core Libraries](#)
[Governing Board](#)
[HotSpot](#)

JEP draft: Epsilon GC: The Arbitrarily Low Overhead Garbage (Non-)Collector

<i>Owner</i>	Aleksey Shipilev
<i>Created</i>	2017/02/14 08:23
<i>Updated</i>	2017/04/01 16:29
<i>Type</i>	Feature
<i>Status</i>	Draft
<i>Component</i>	hotspot / gc
<i>Scope</i>	Implementation
<i>Effort</i>	S
<i>Duration</i>	S
<i>Priority</i>	4
<i>Issue</i>	8174901


Summary

Develop a GC that only handles memory allocation, but does not implement any actual memory reclamation mechanism. Once available Java heap is exhausted, perform the orderly JVM shutdown.

Evaluating cost of GC

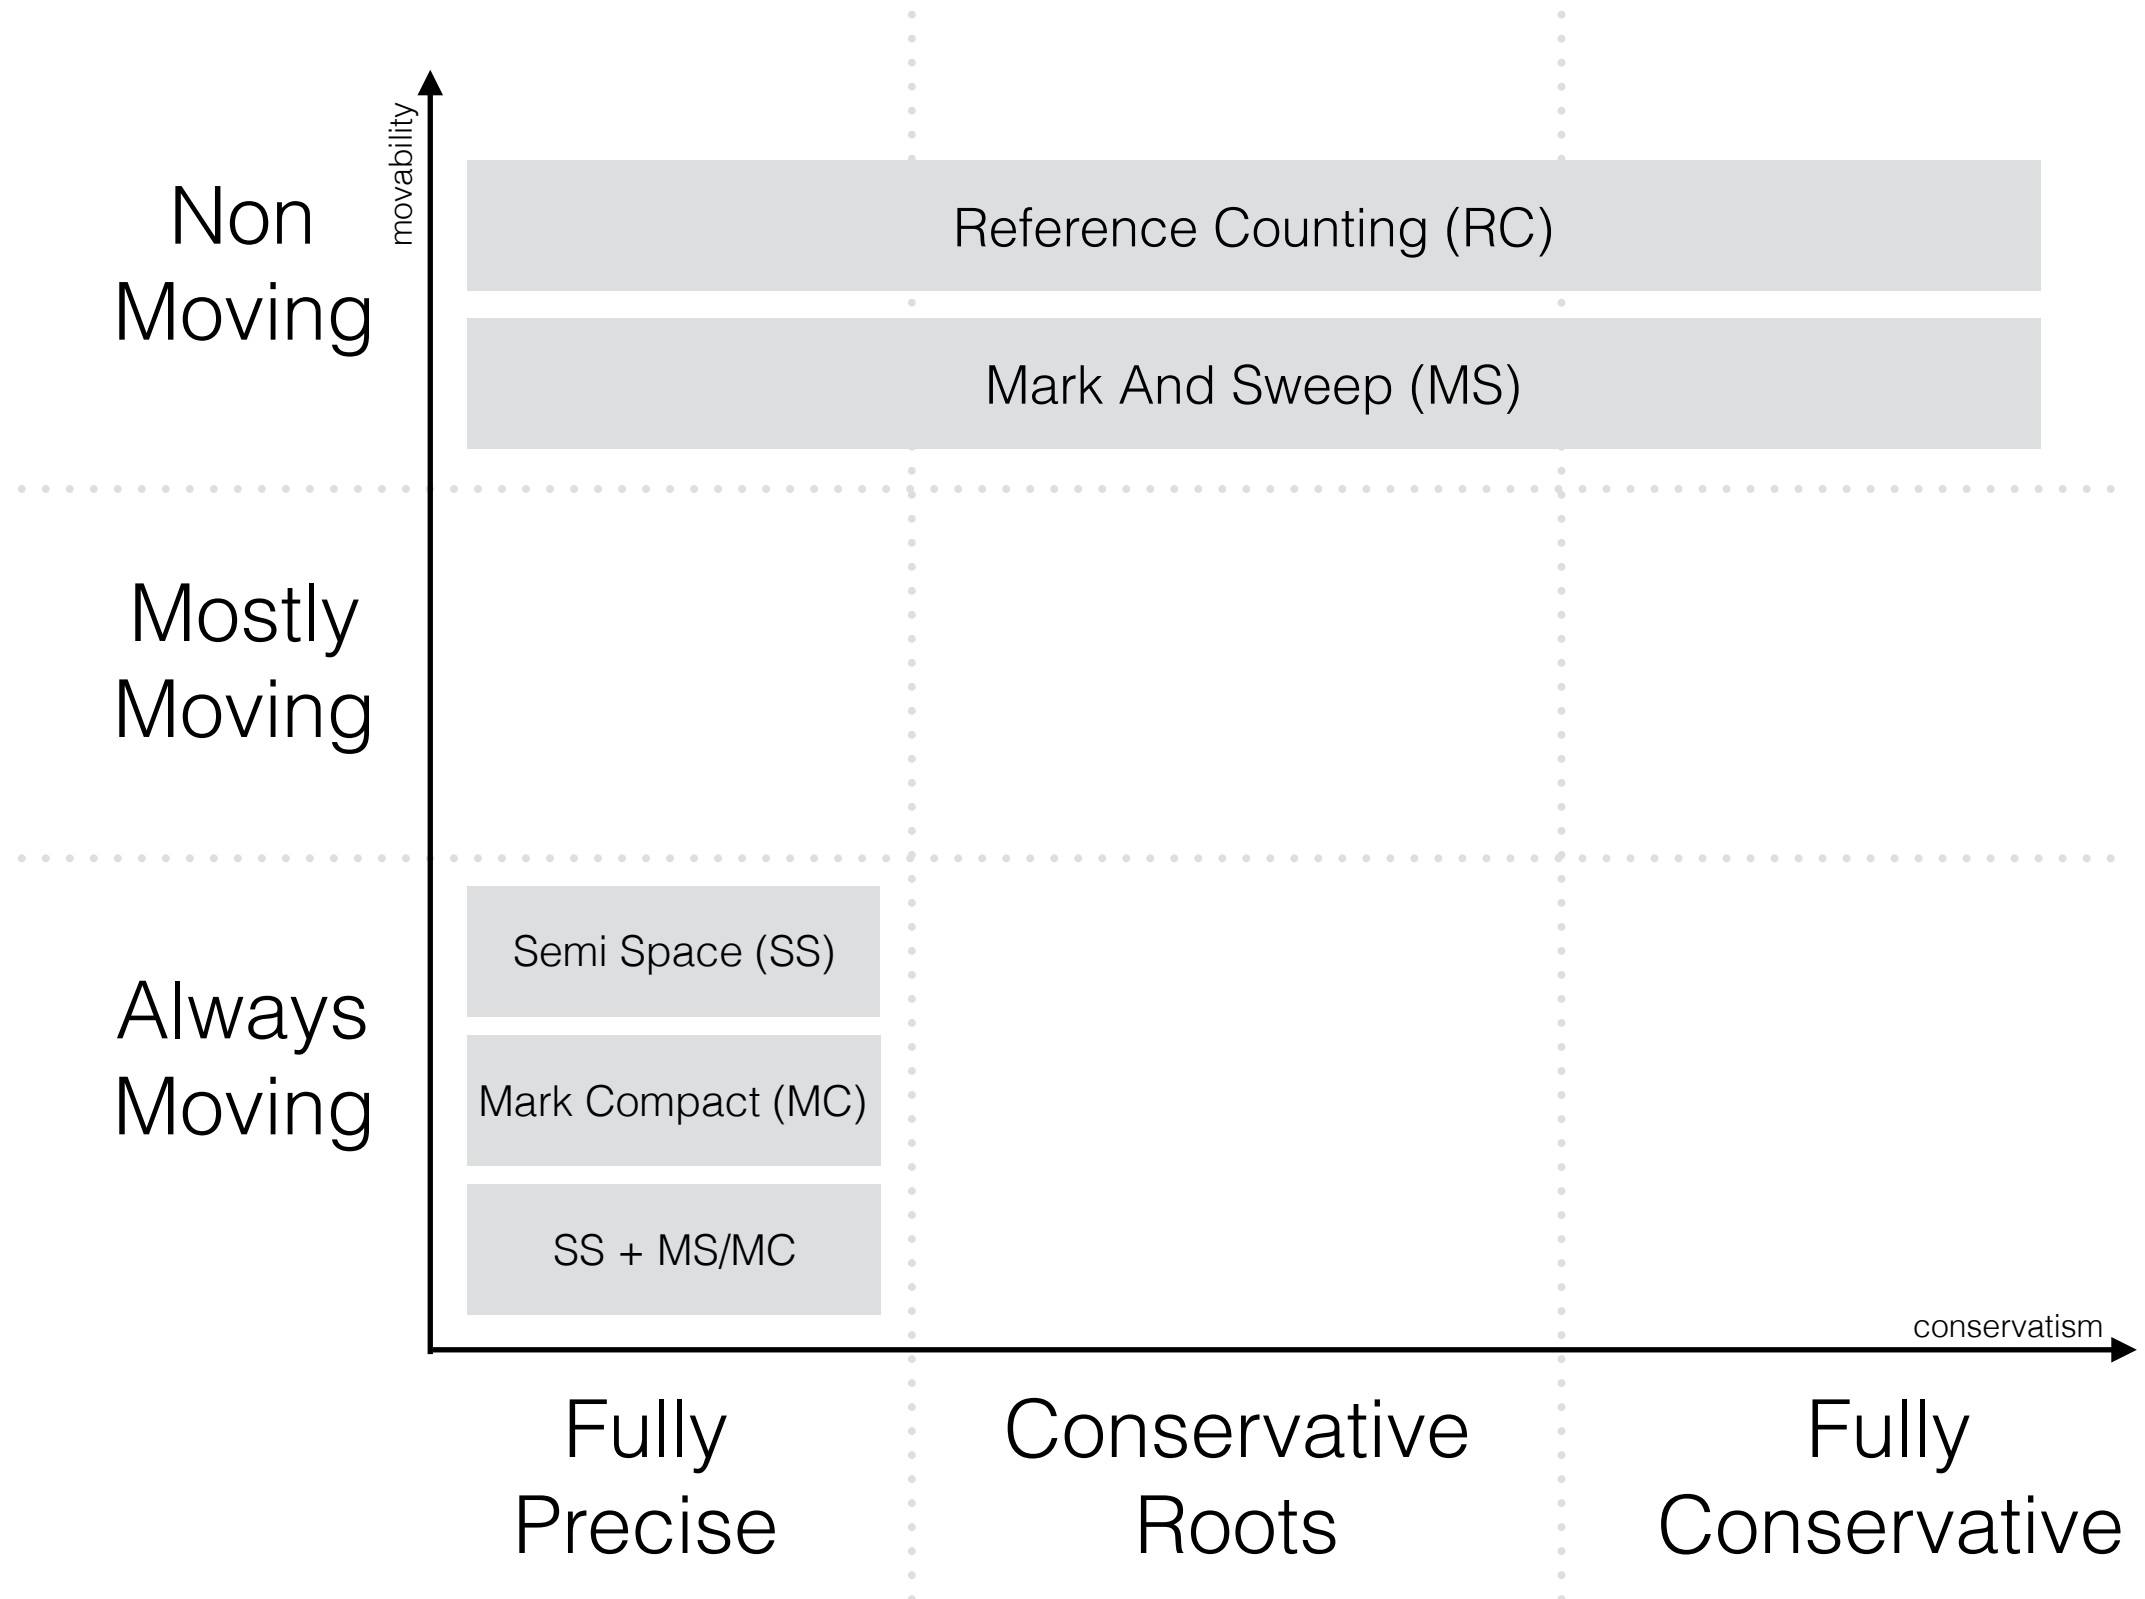
None

Boehm

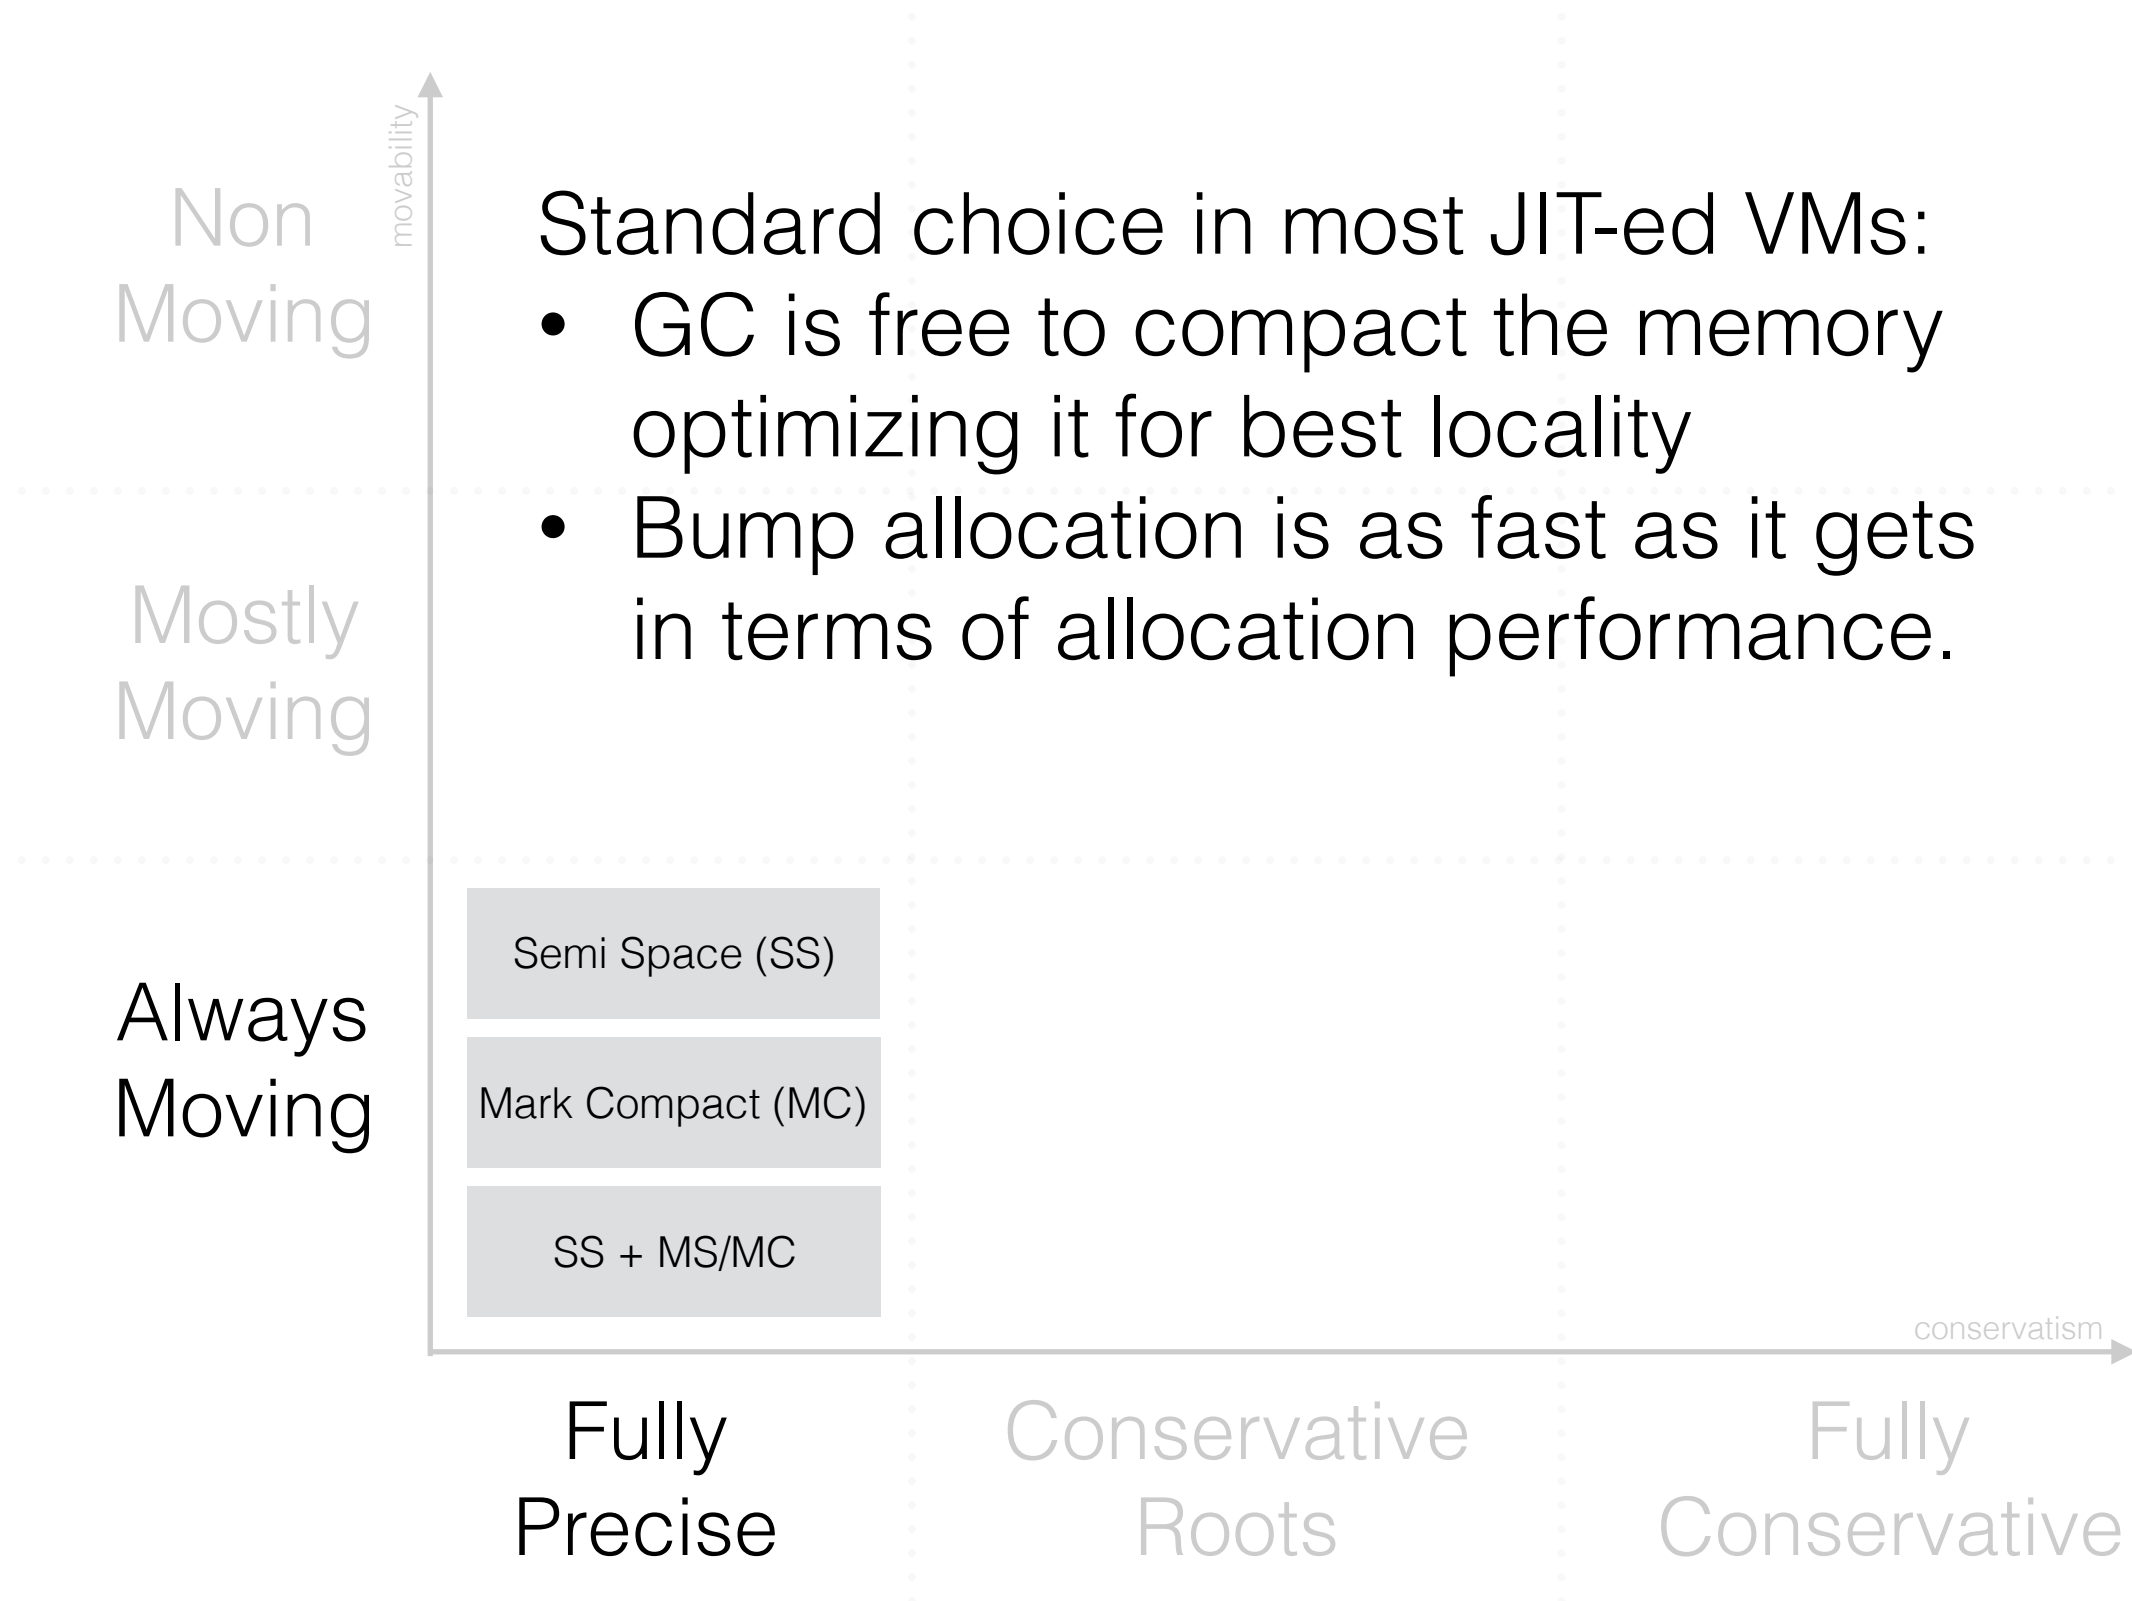


`nativeGC := “?”`


GCs in the wild



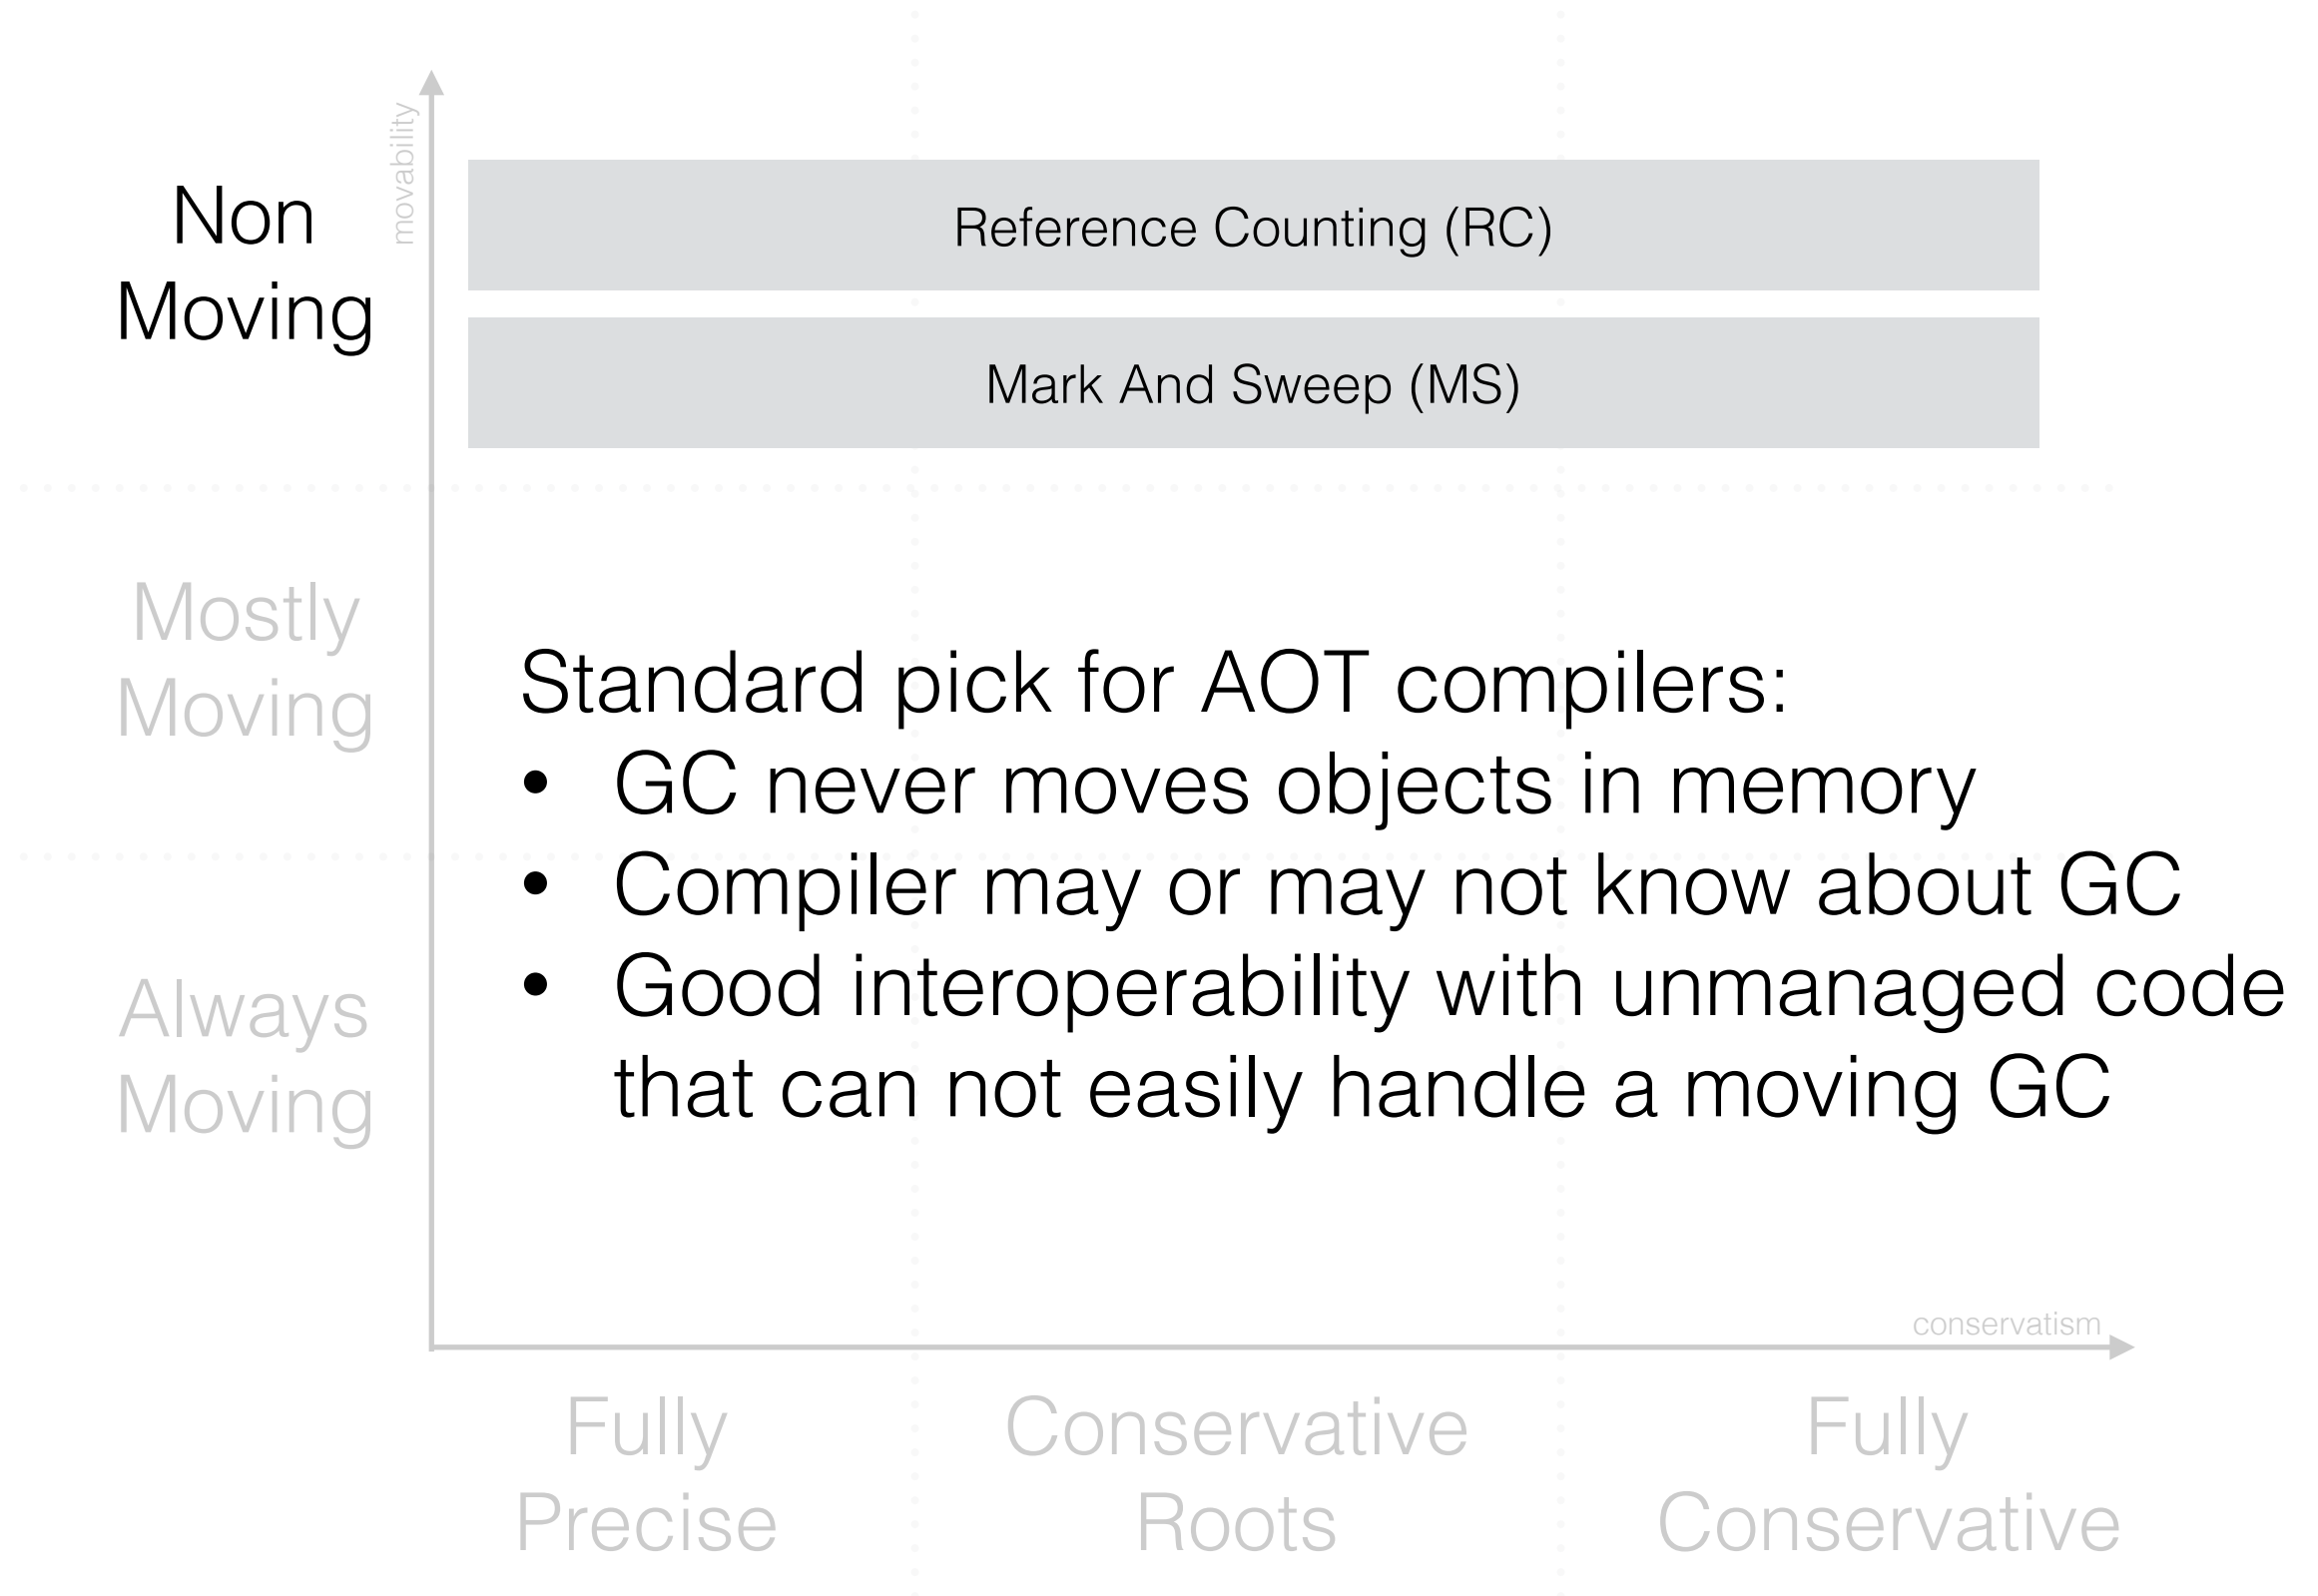
Always Moving GCs



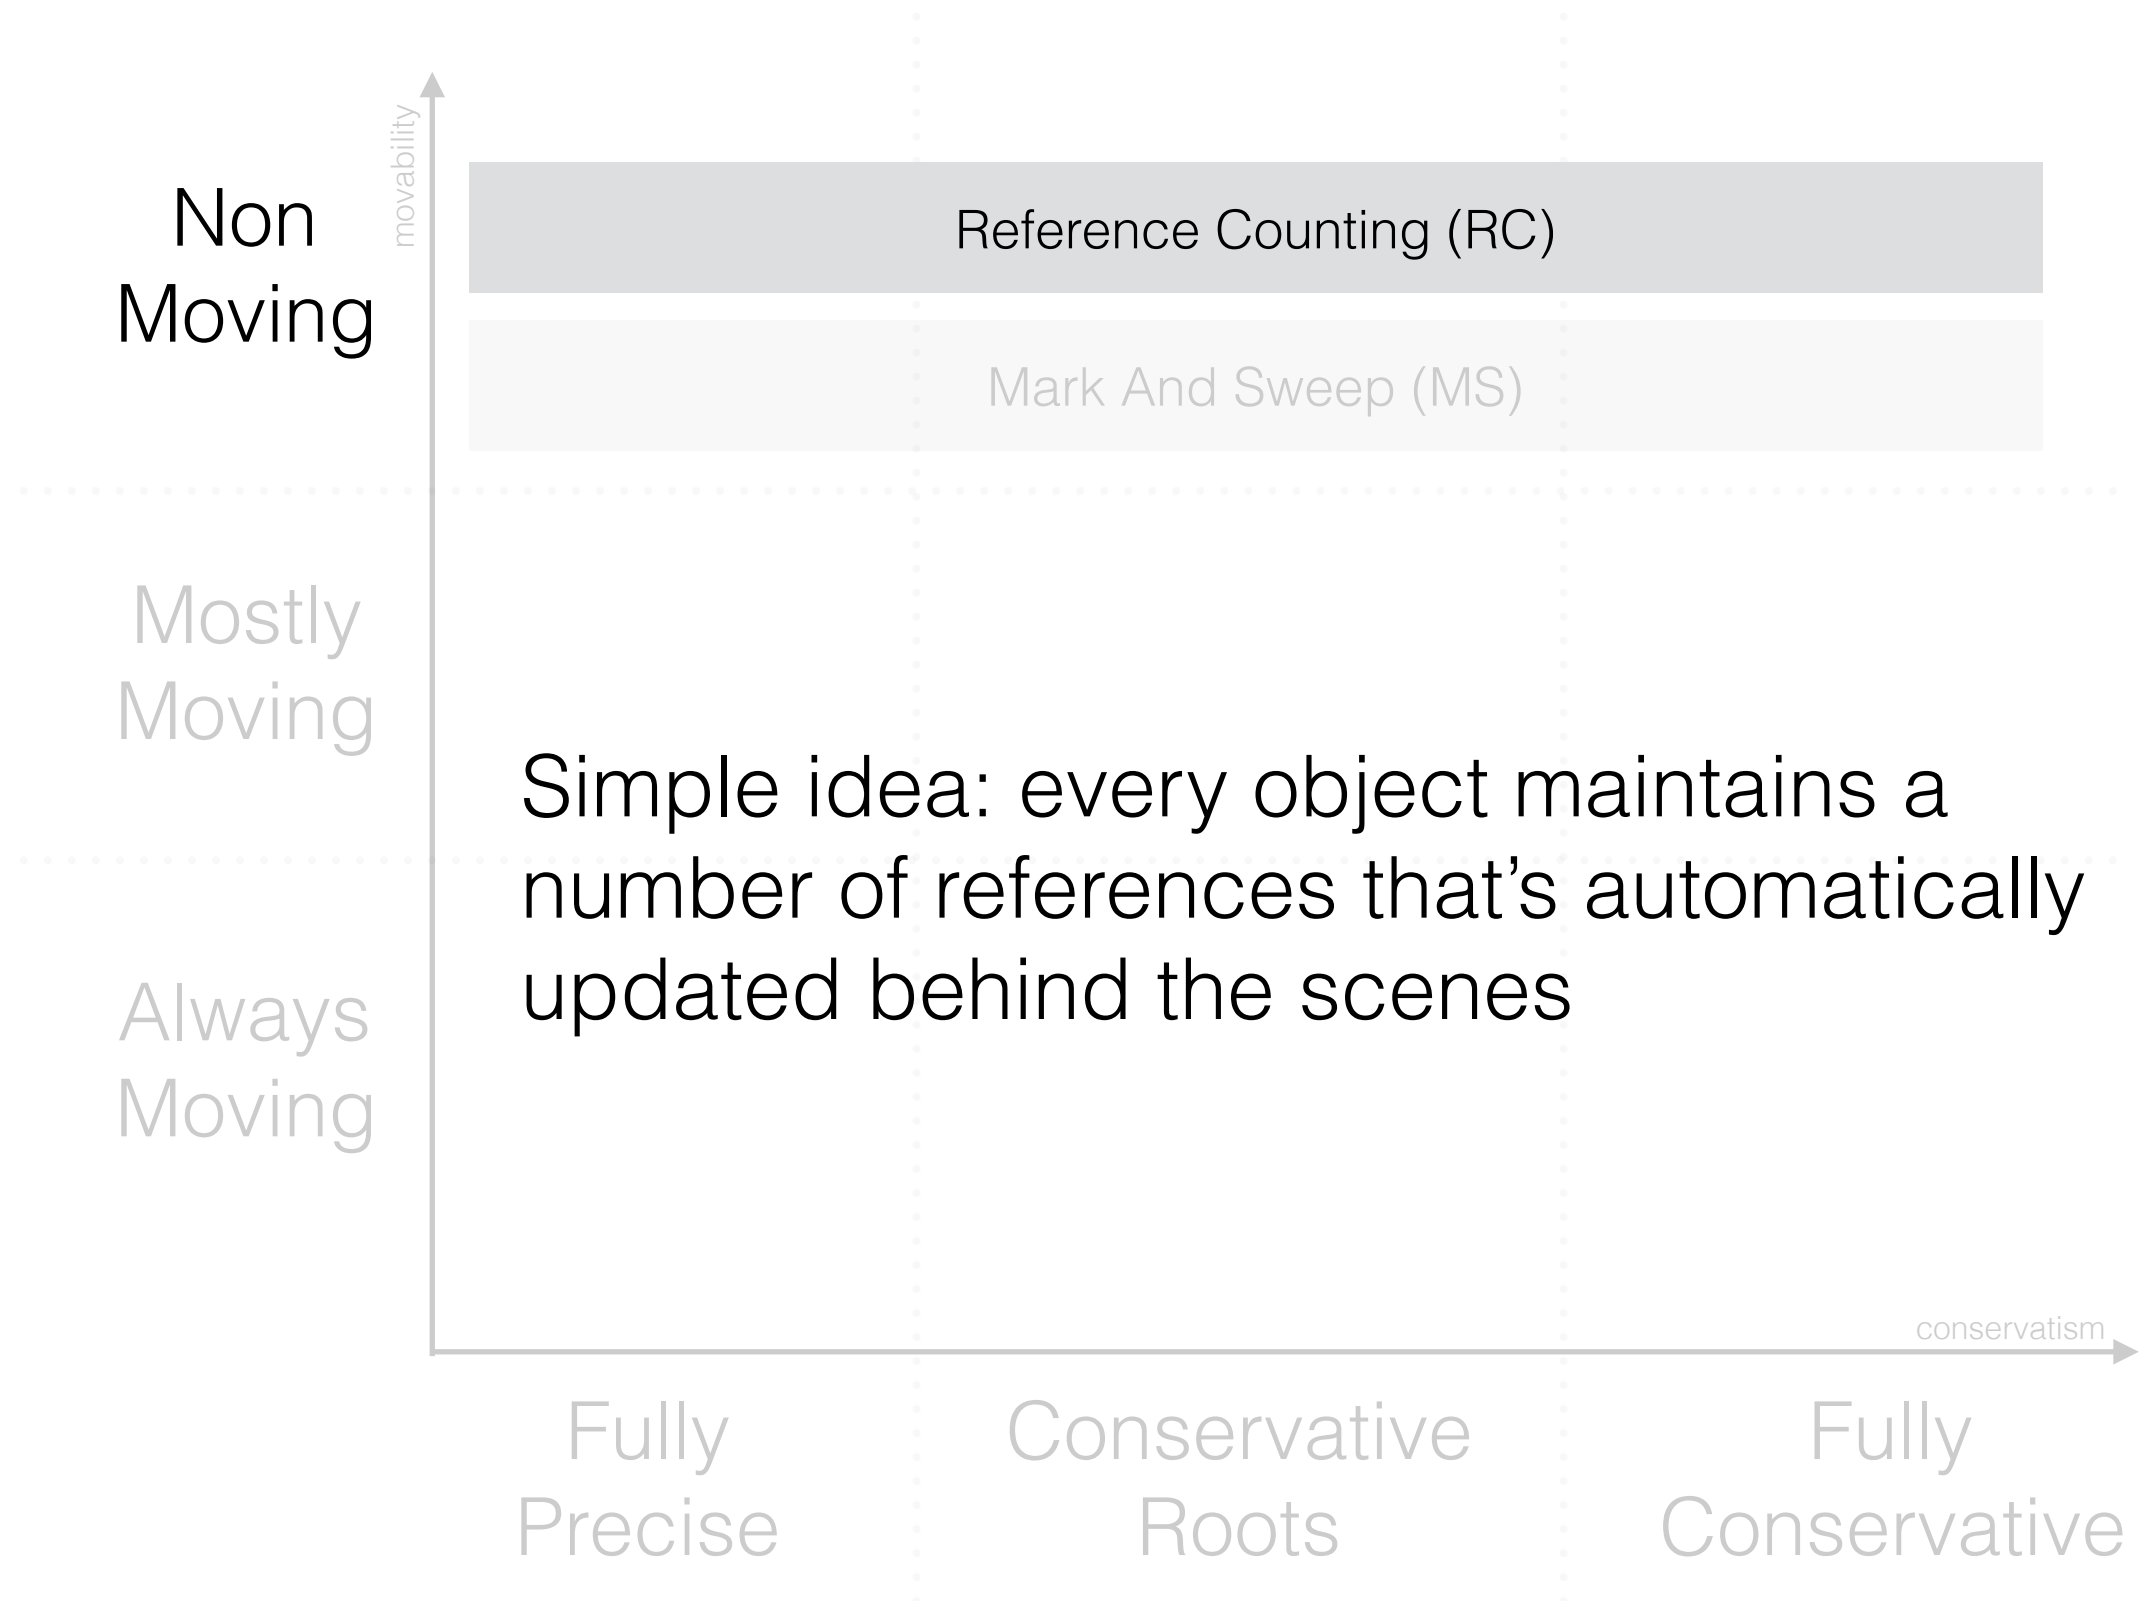
Always Moving GCs



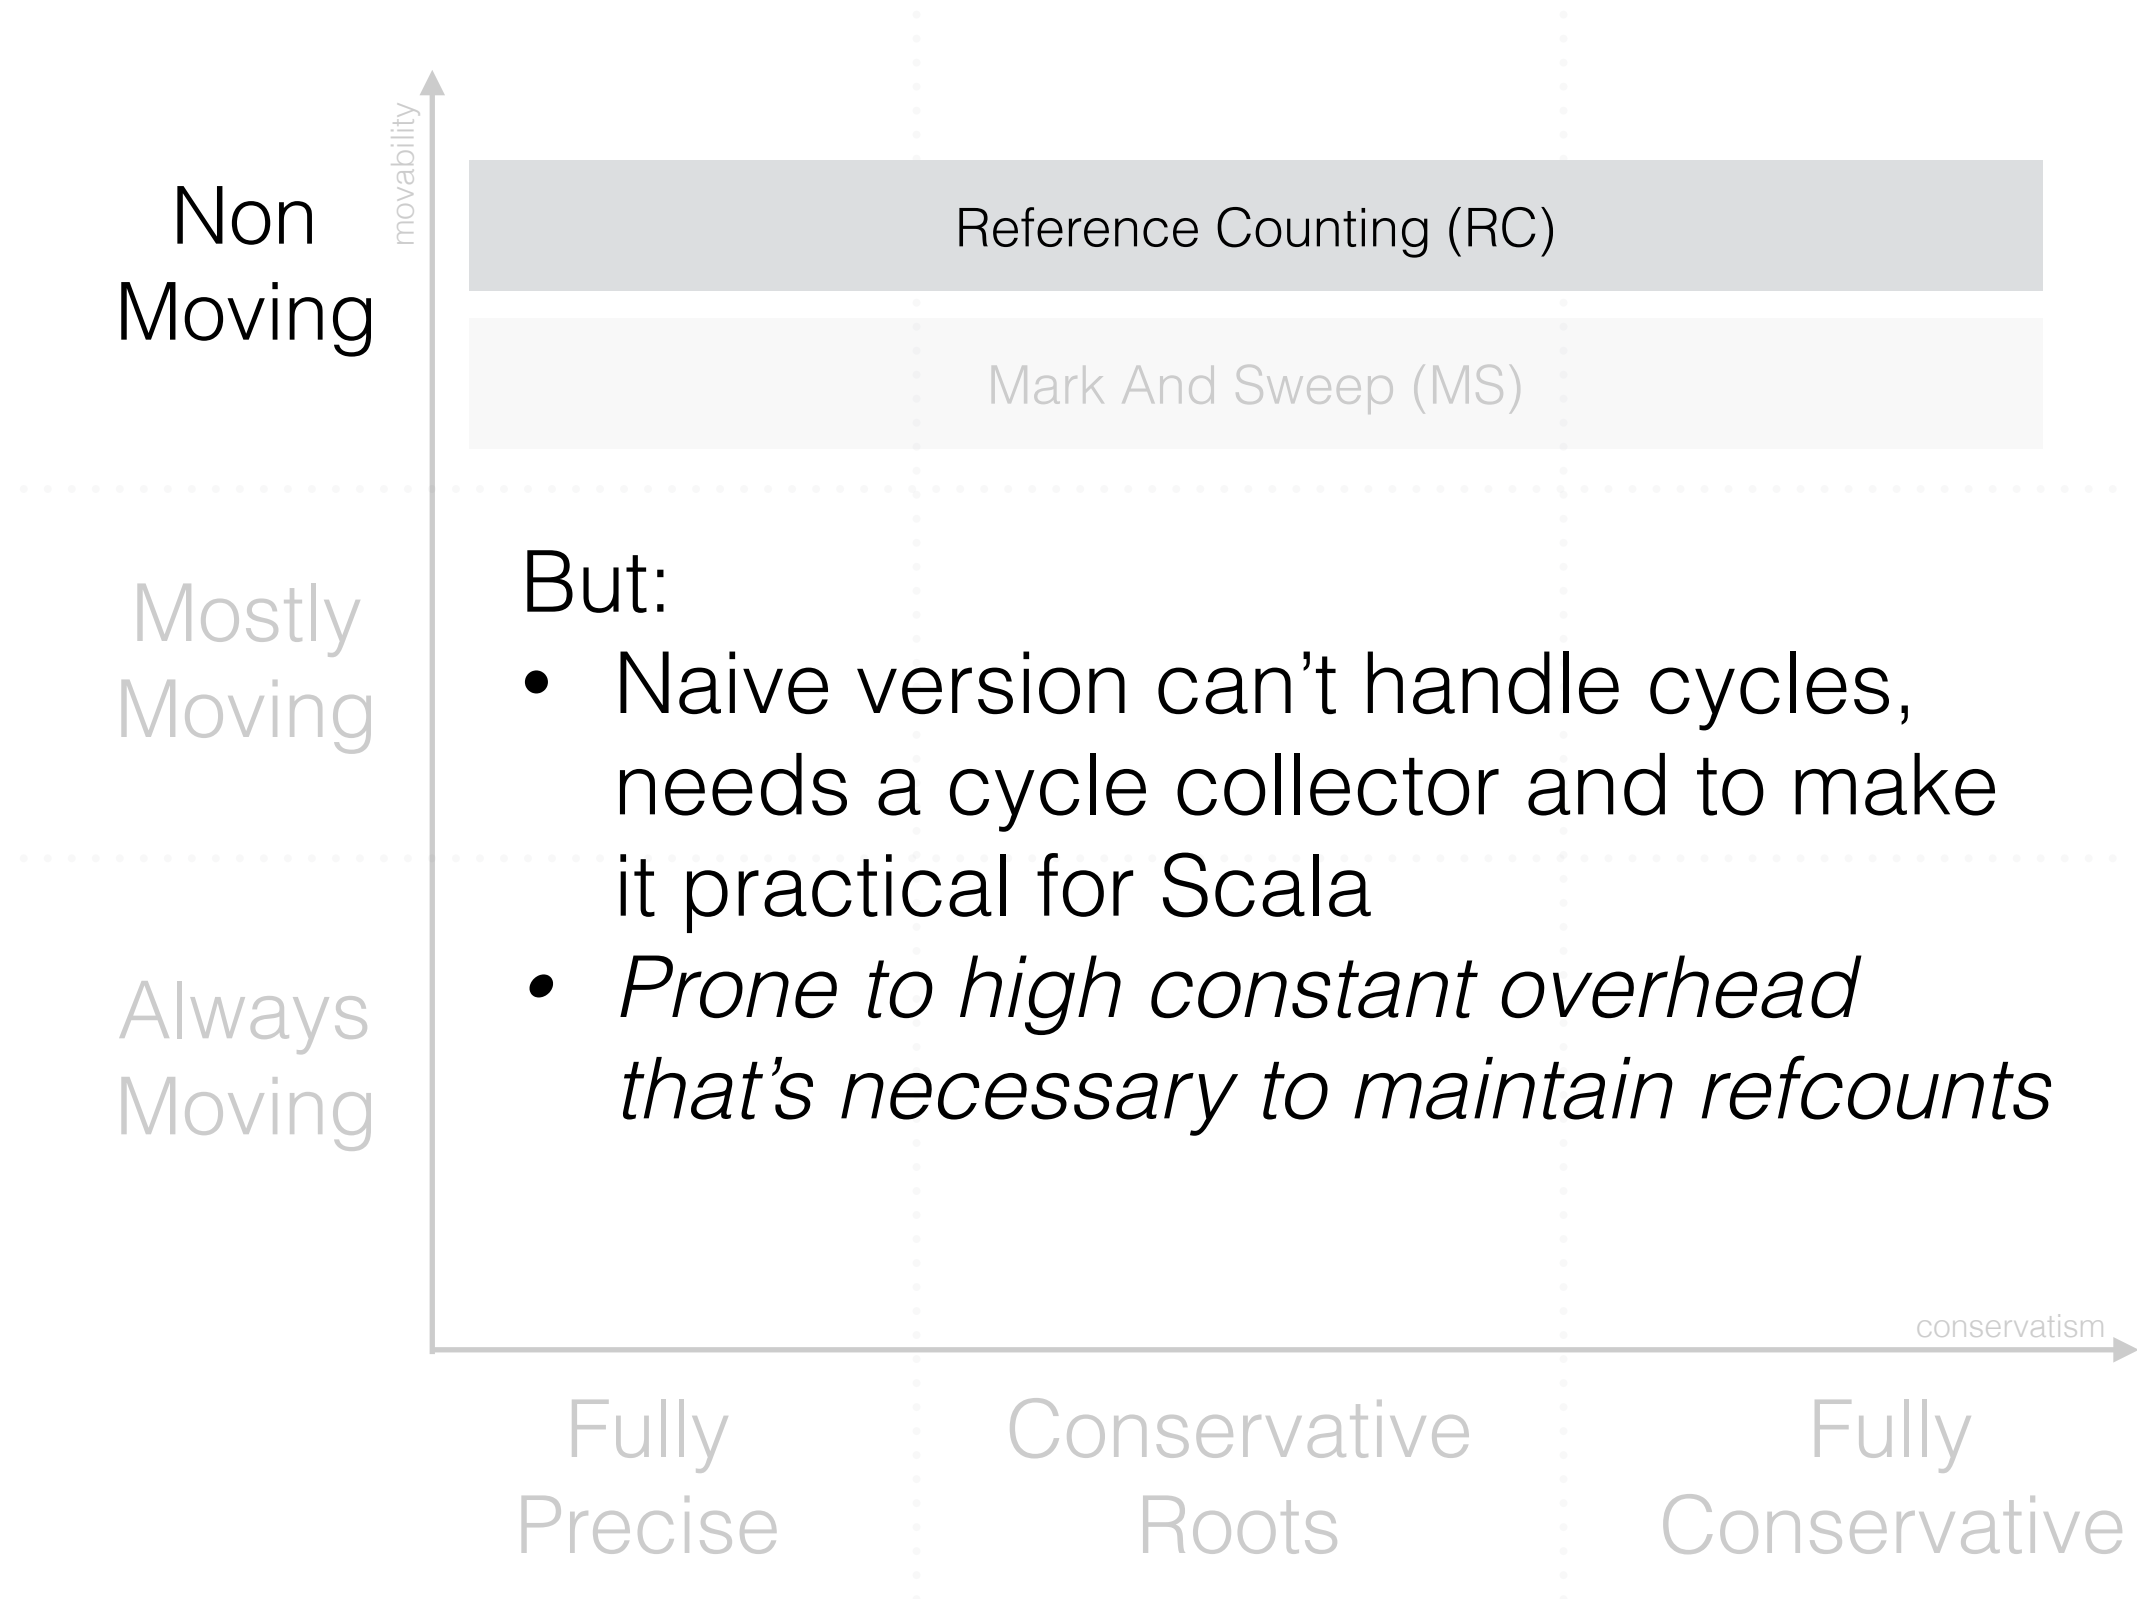
Non-Moving GCs



Reference Counting



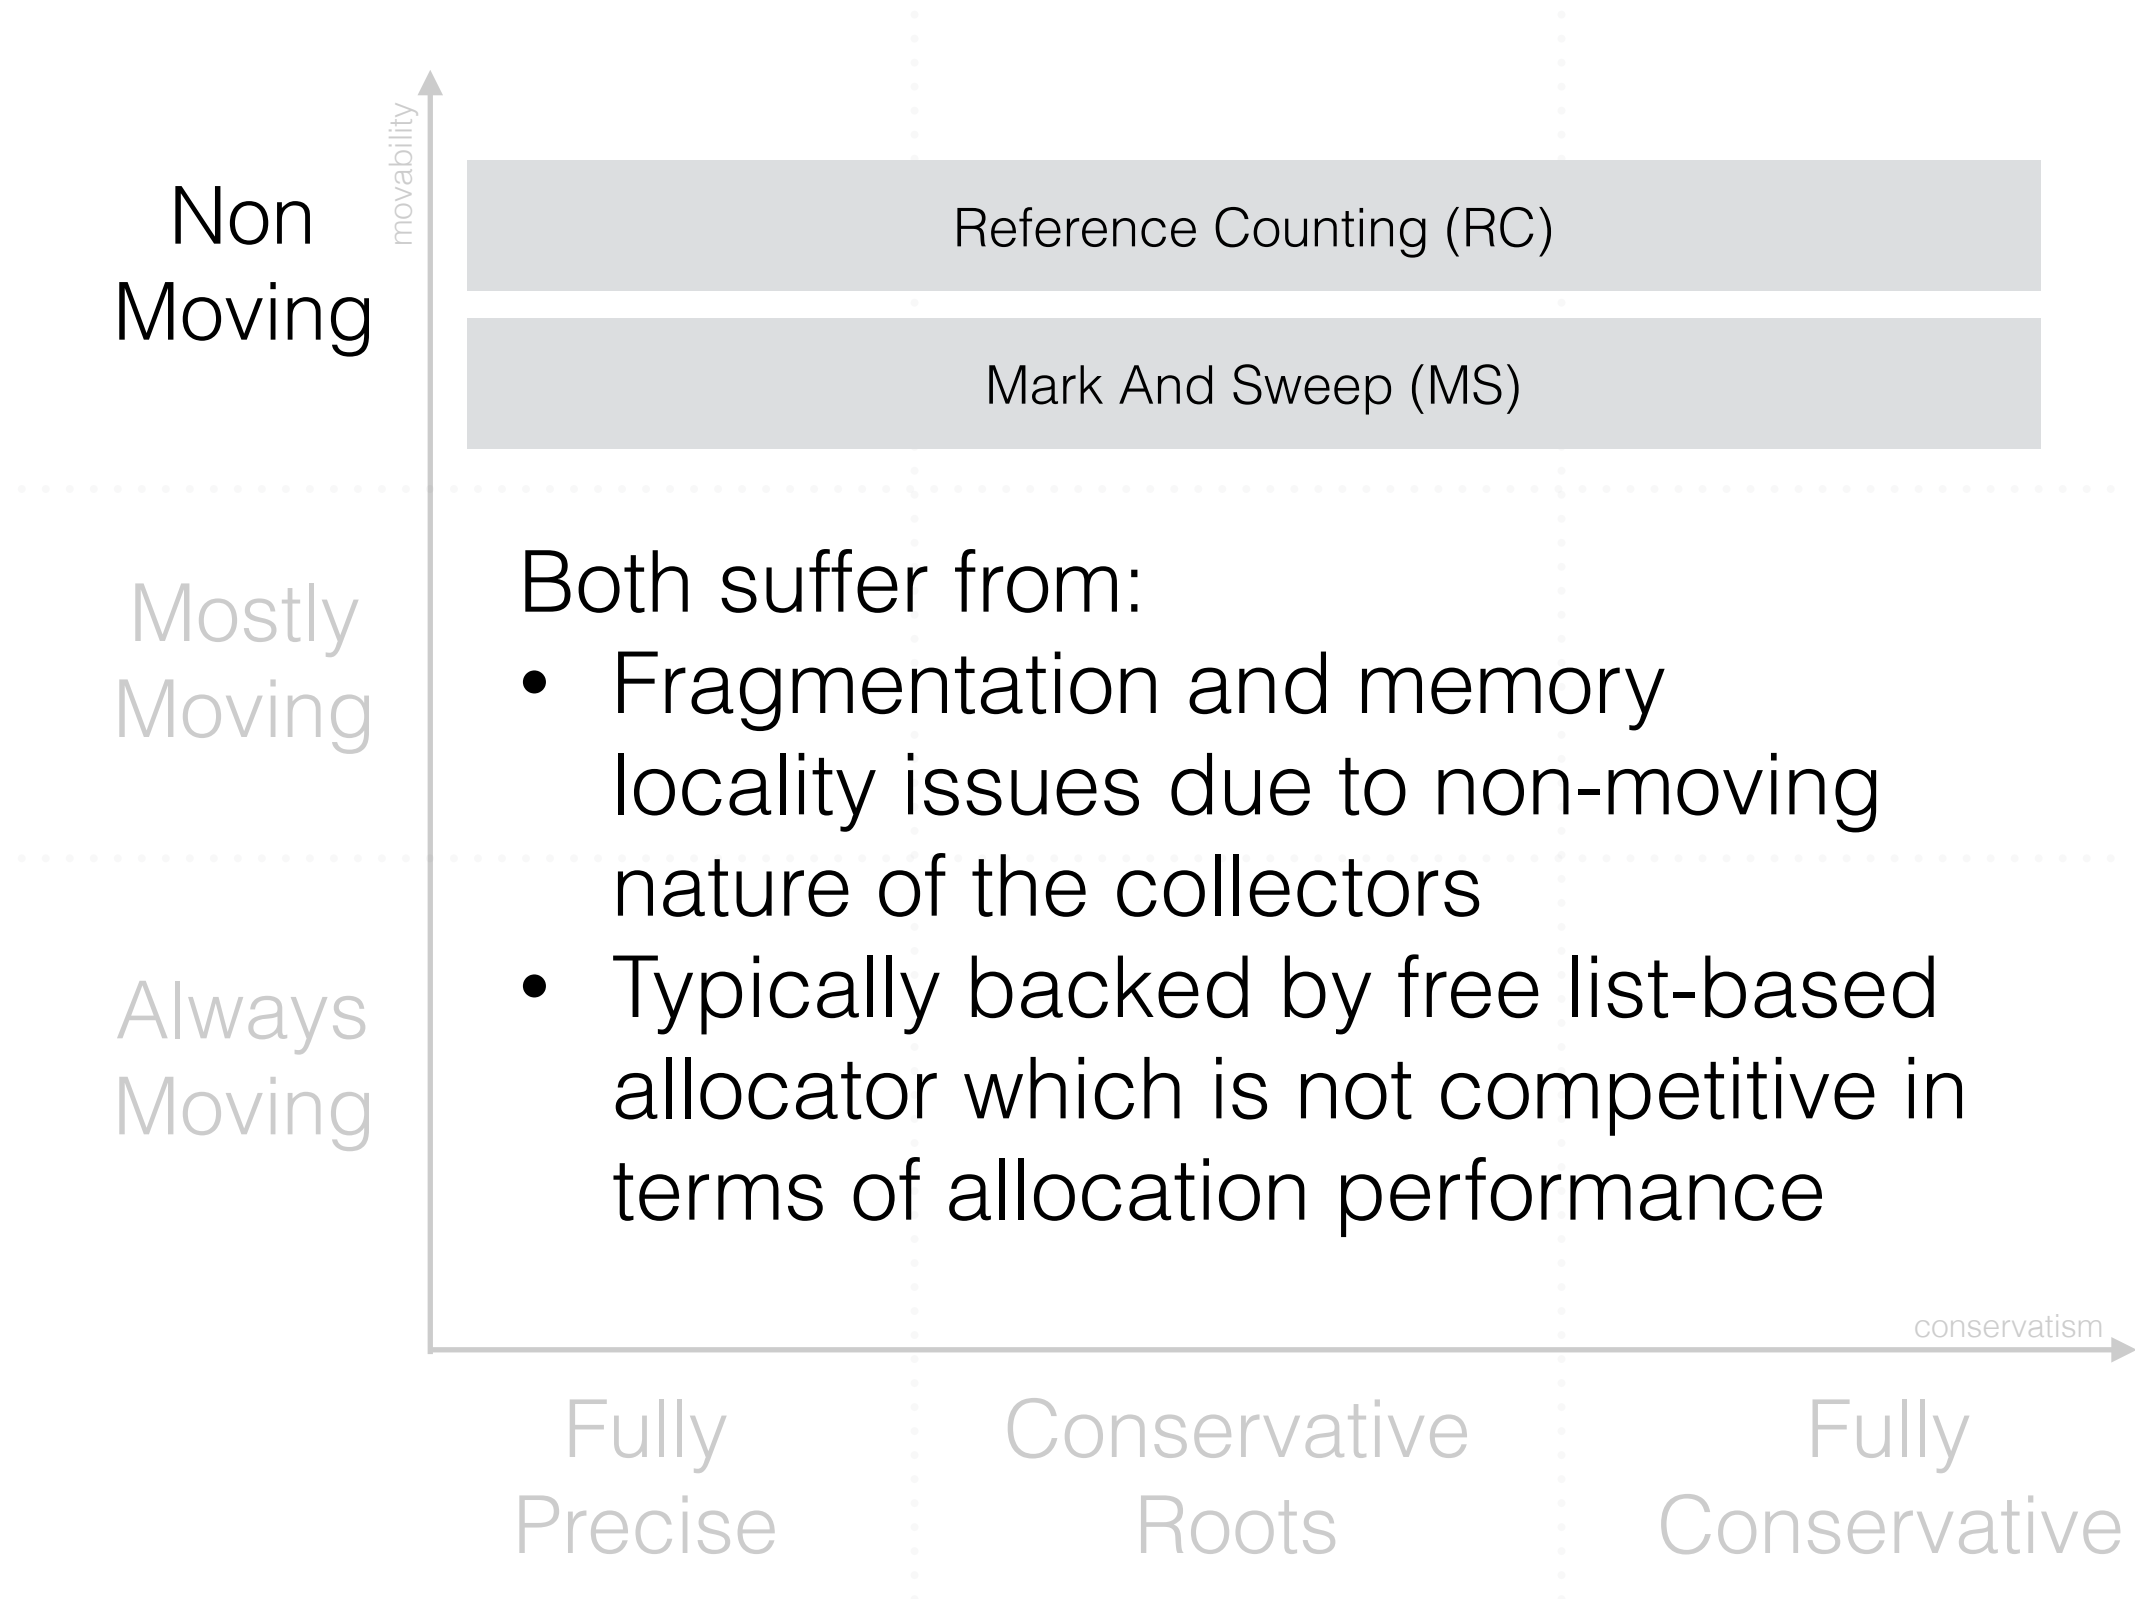
Reference Counting




Reference Counting



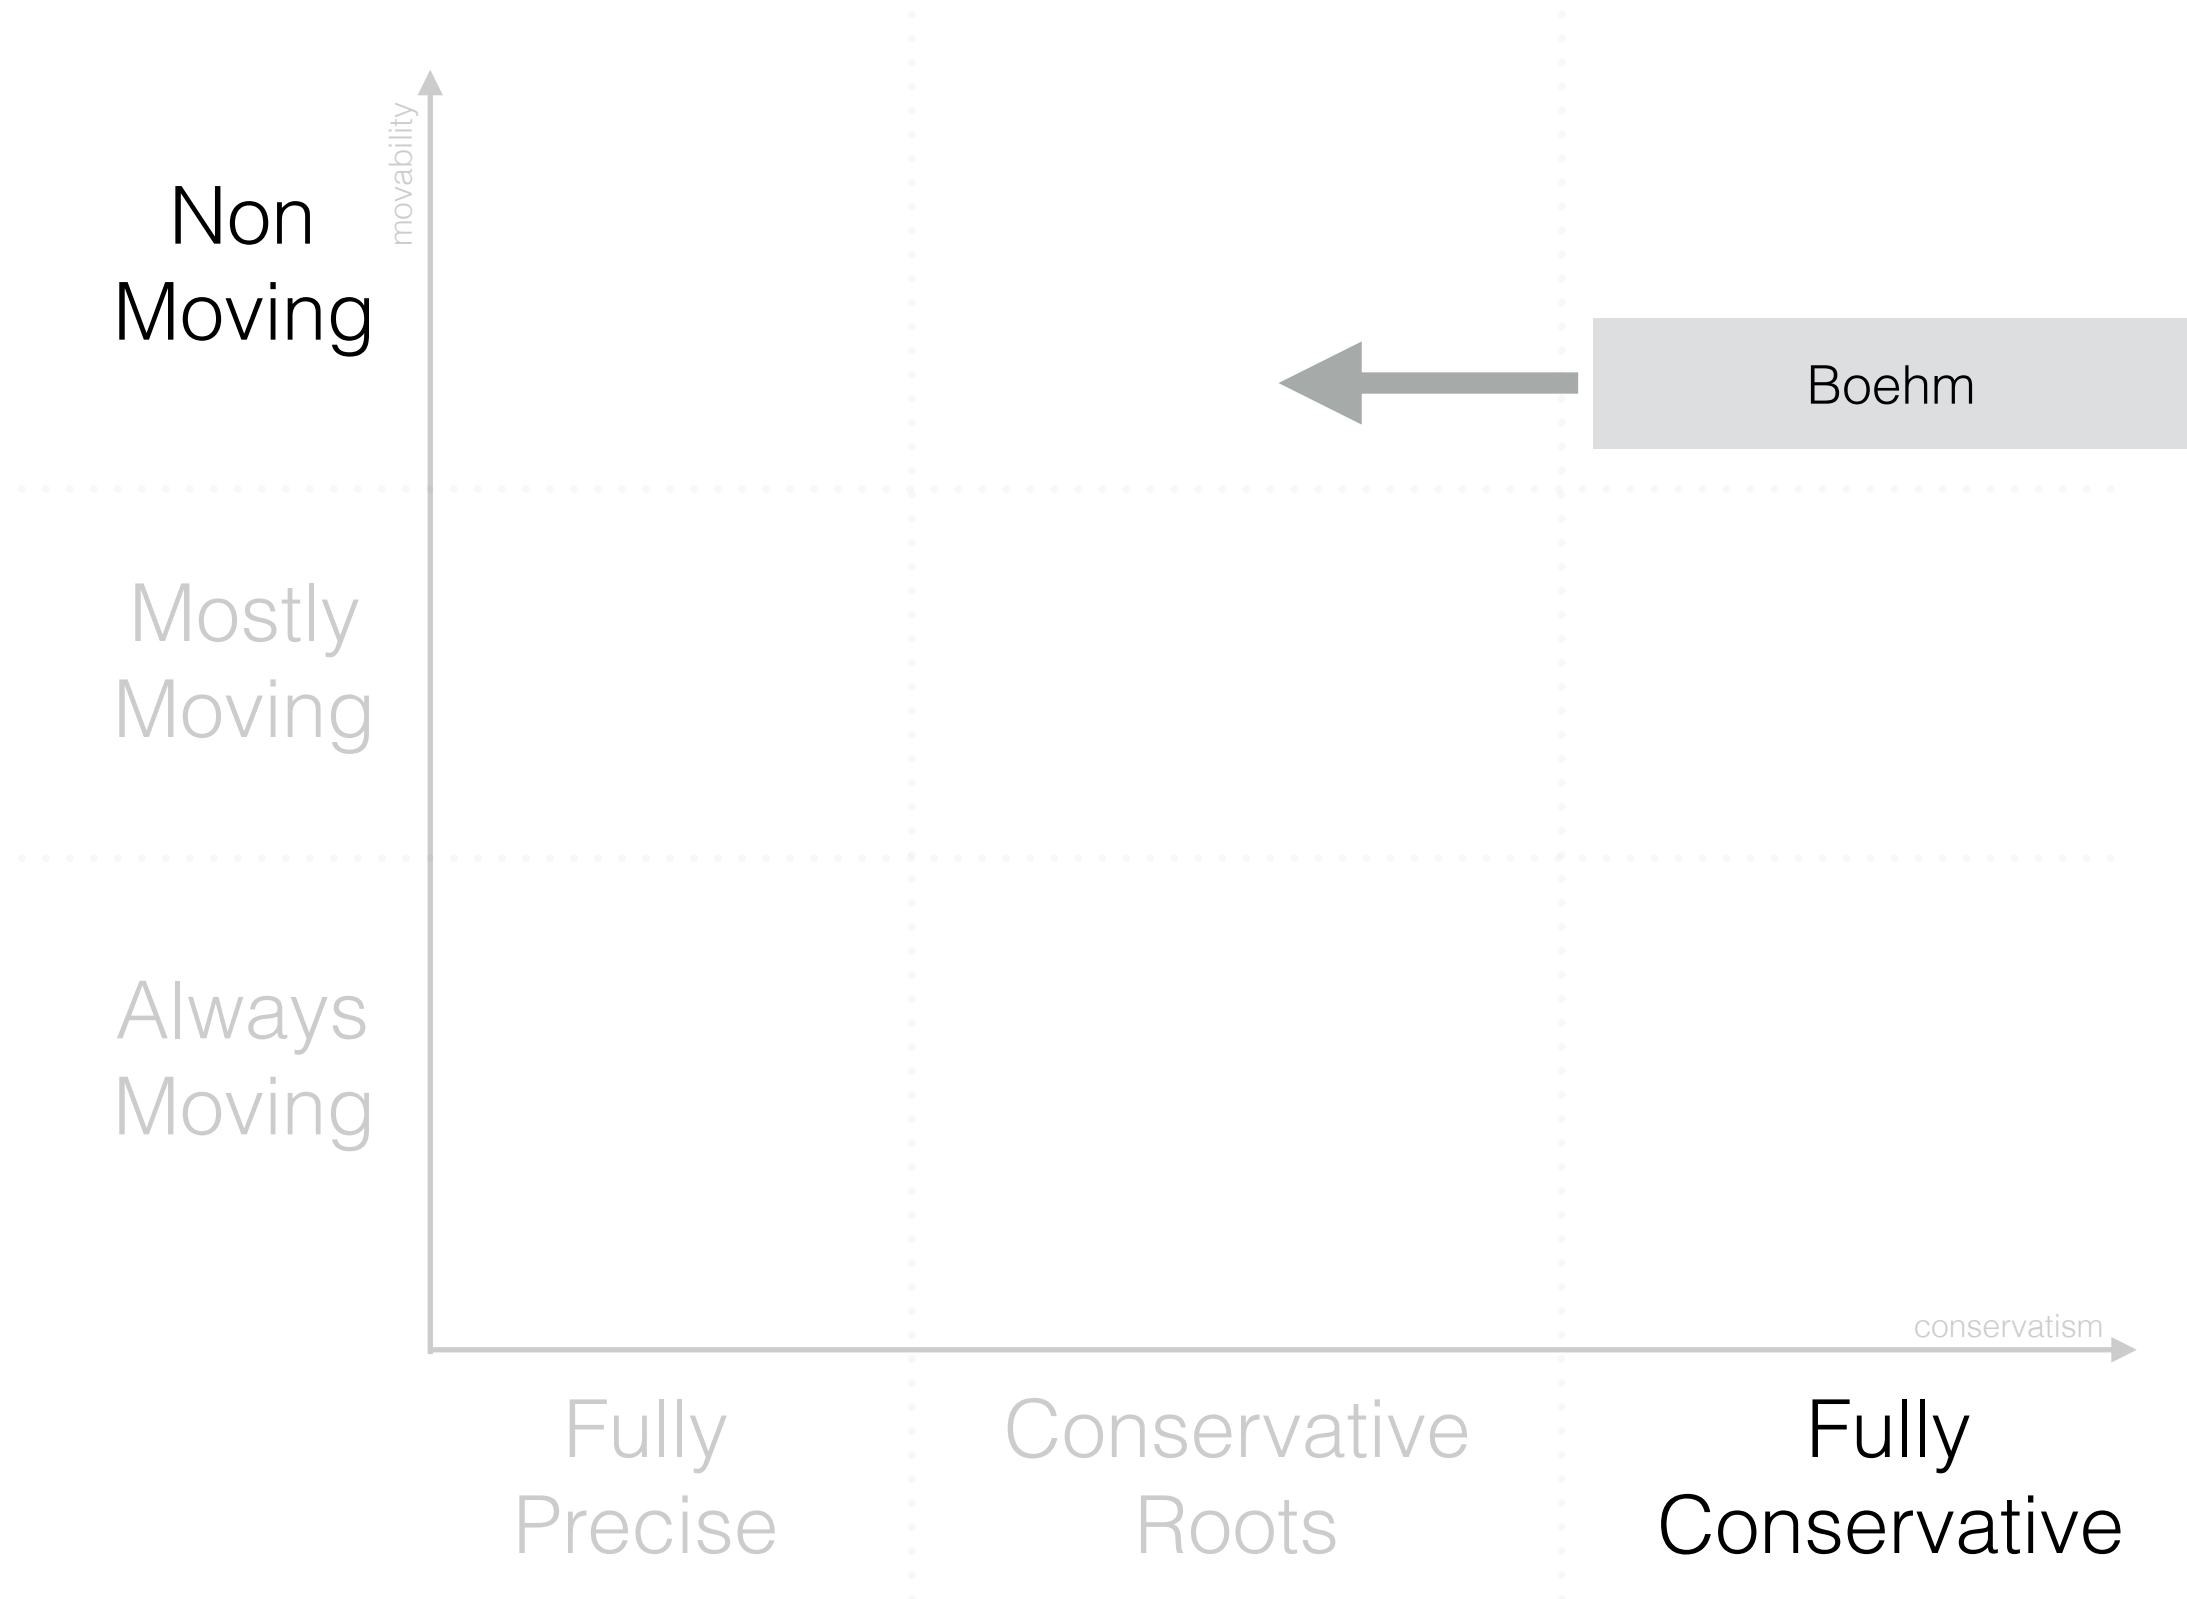
Non Moving GCs



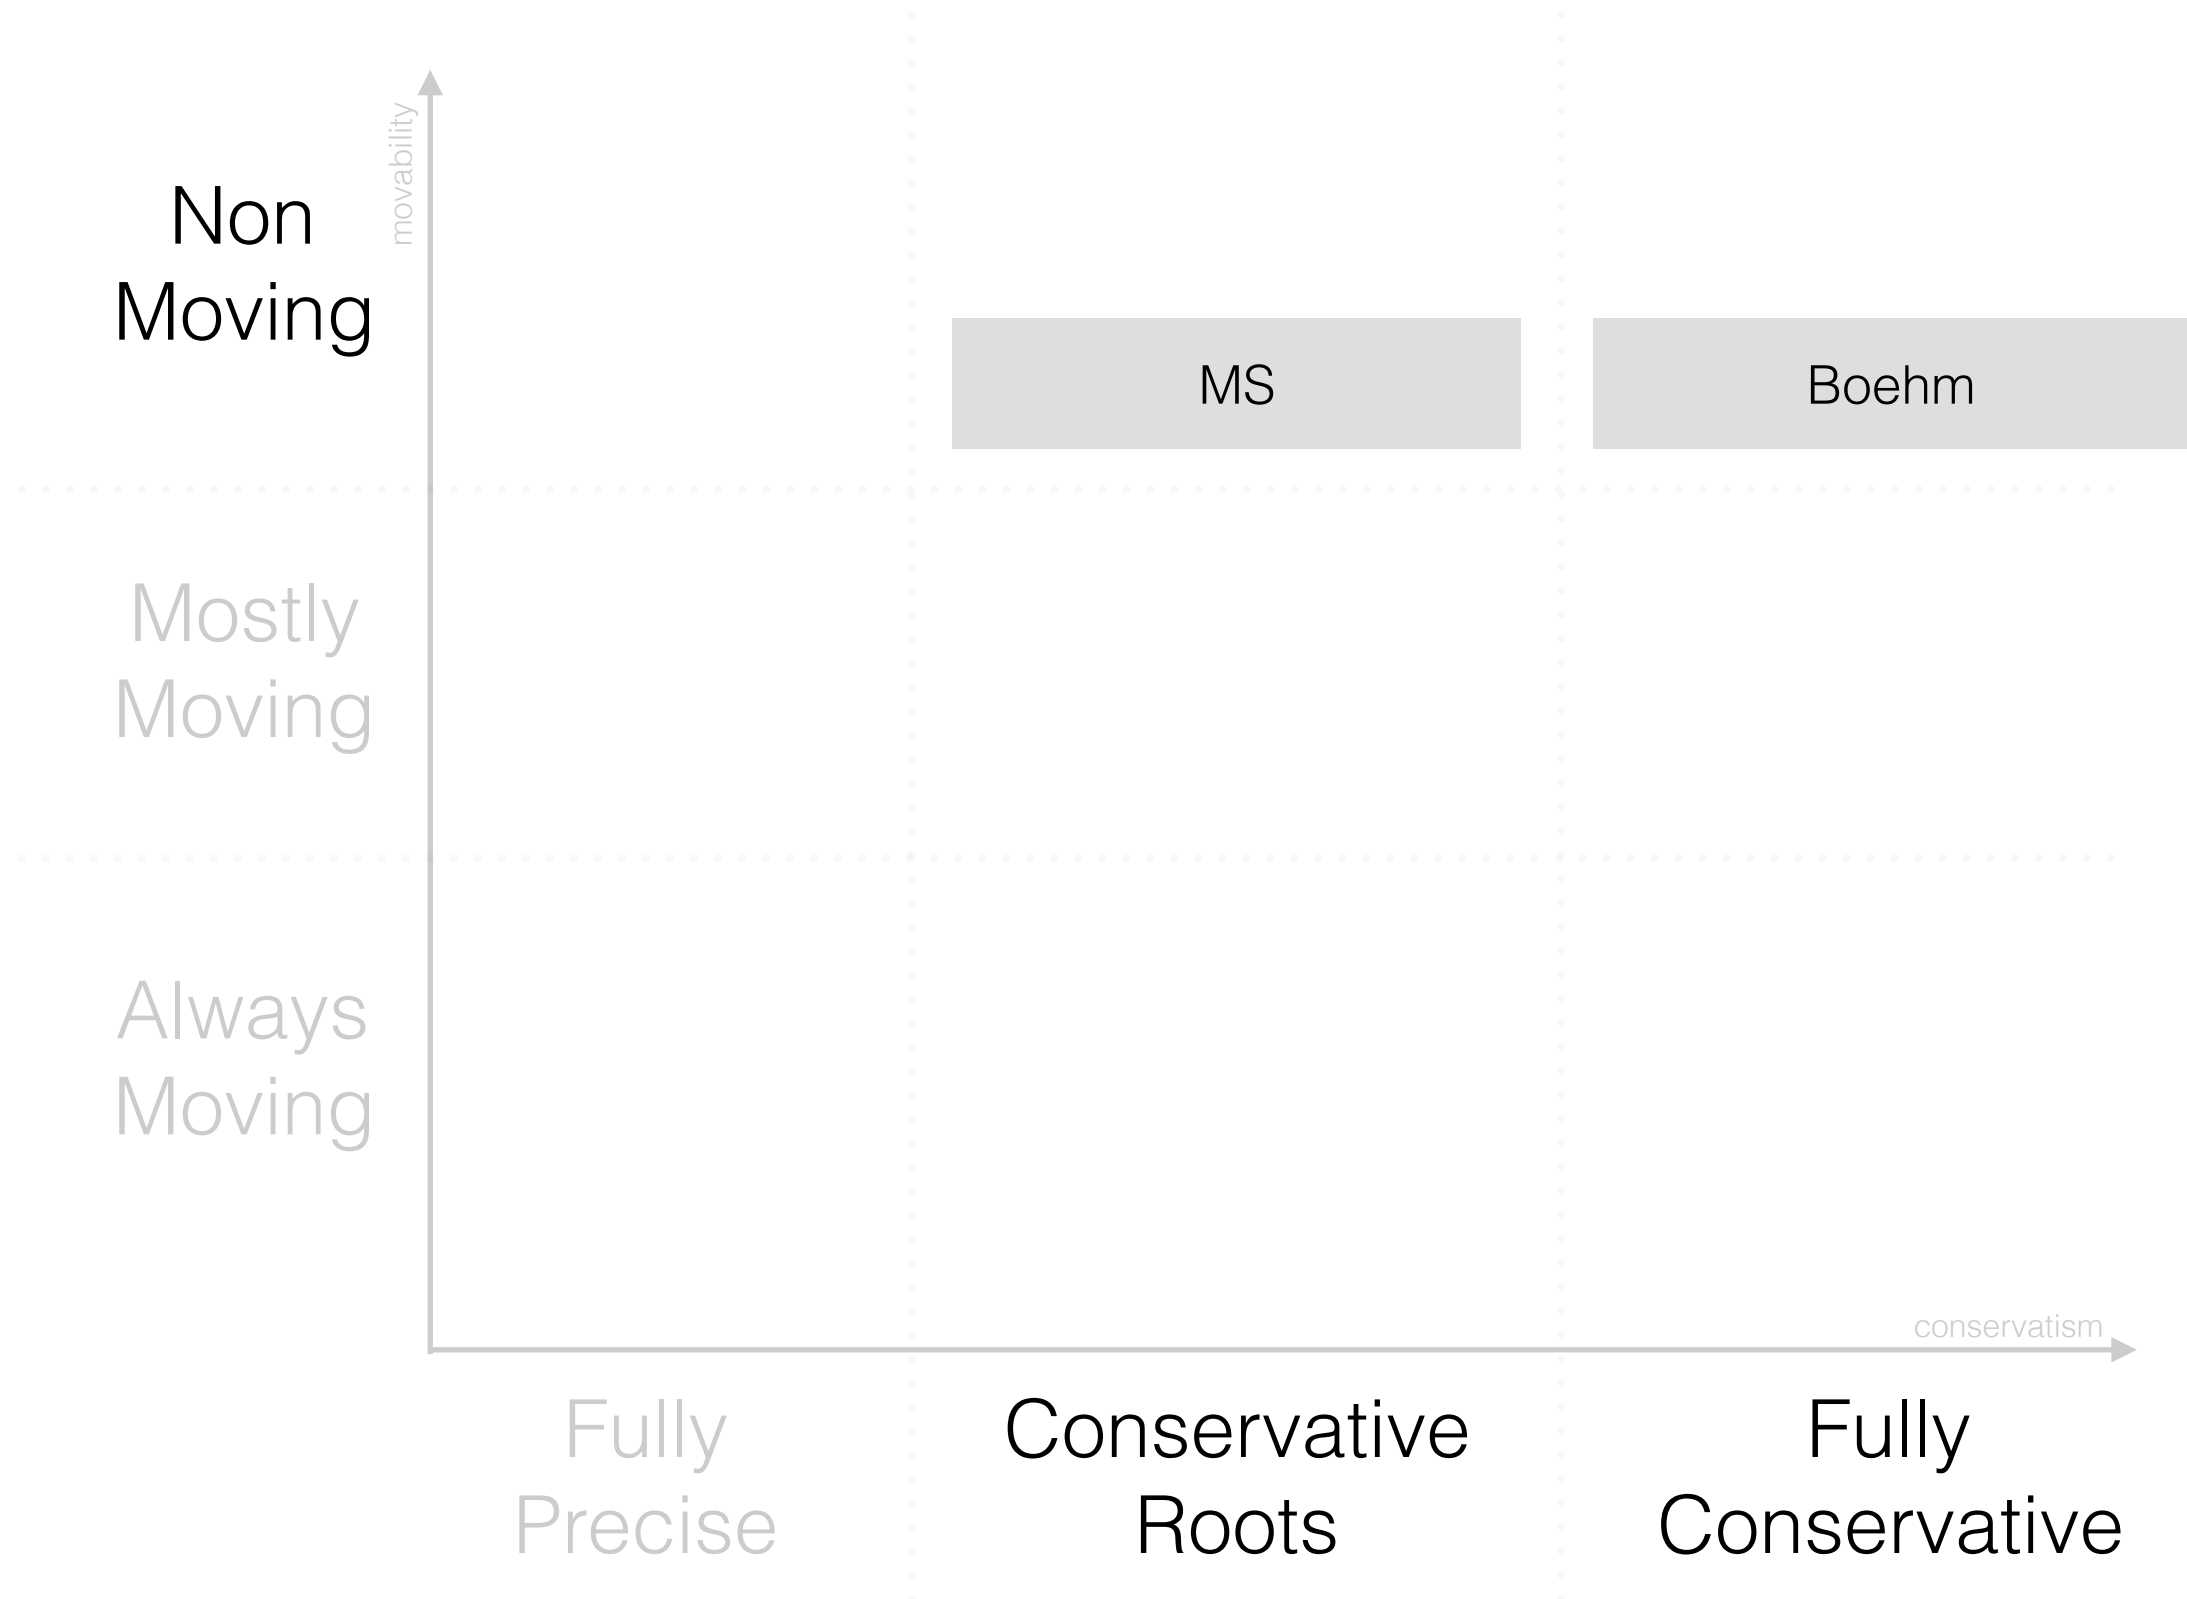
Status quo revisited



Status quo revisited




Initial experiments




MS

Stack







Modules



MS

Stack




Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase

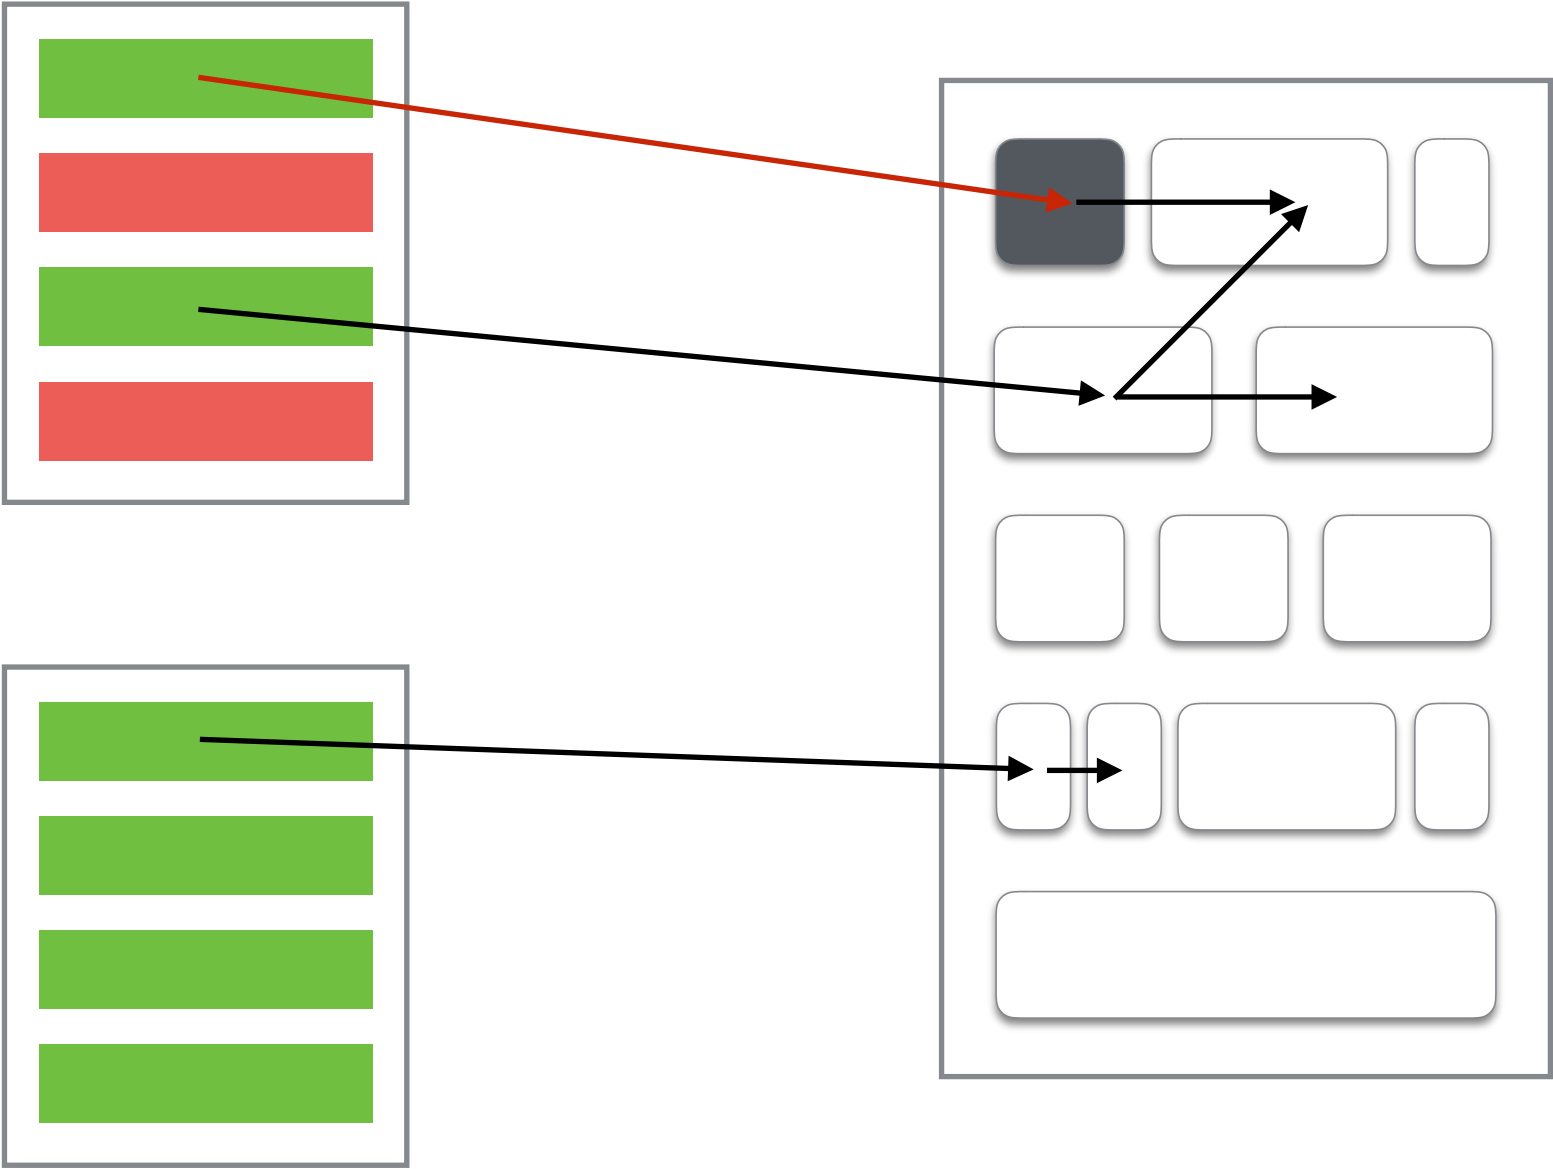


Modules

MS

Stack

Marking Phase

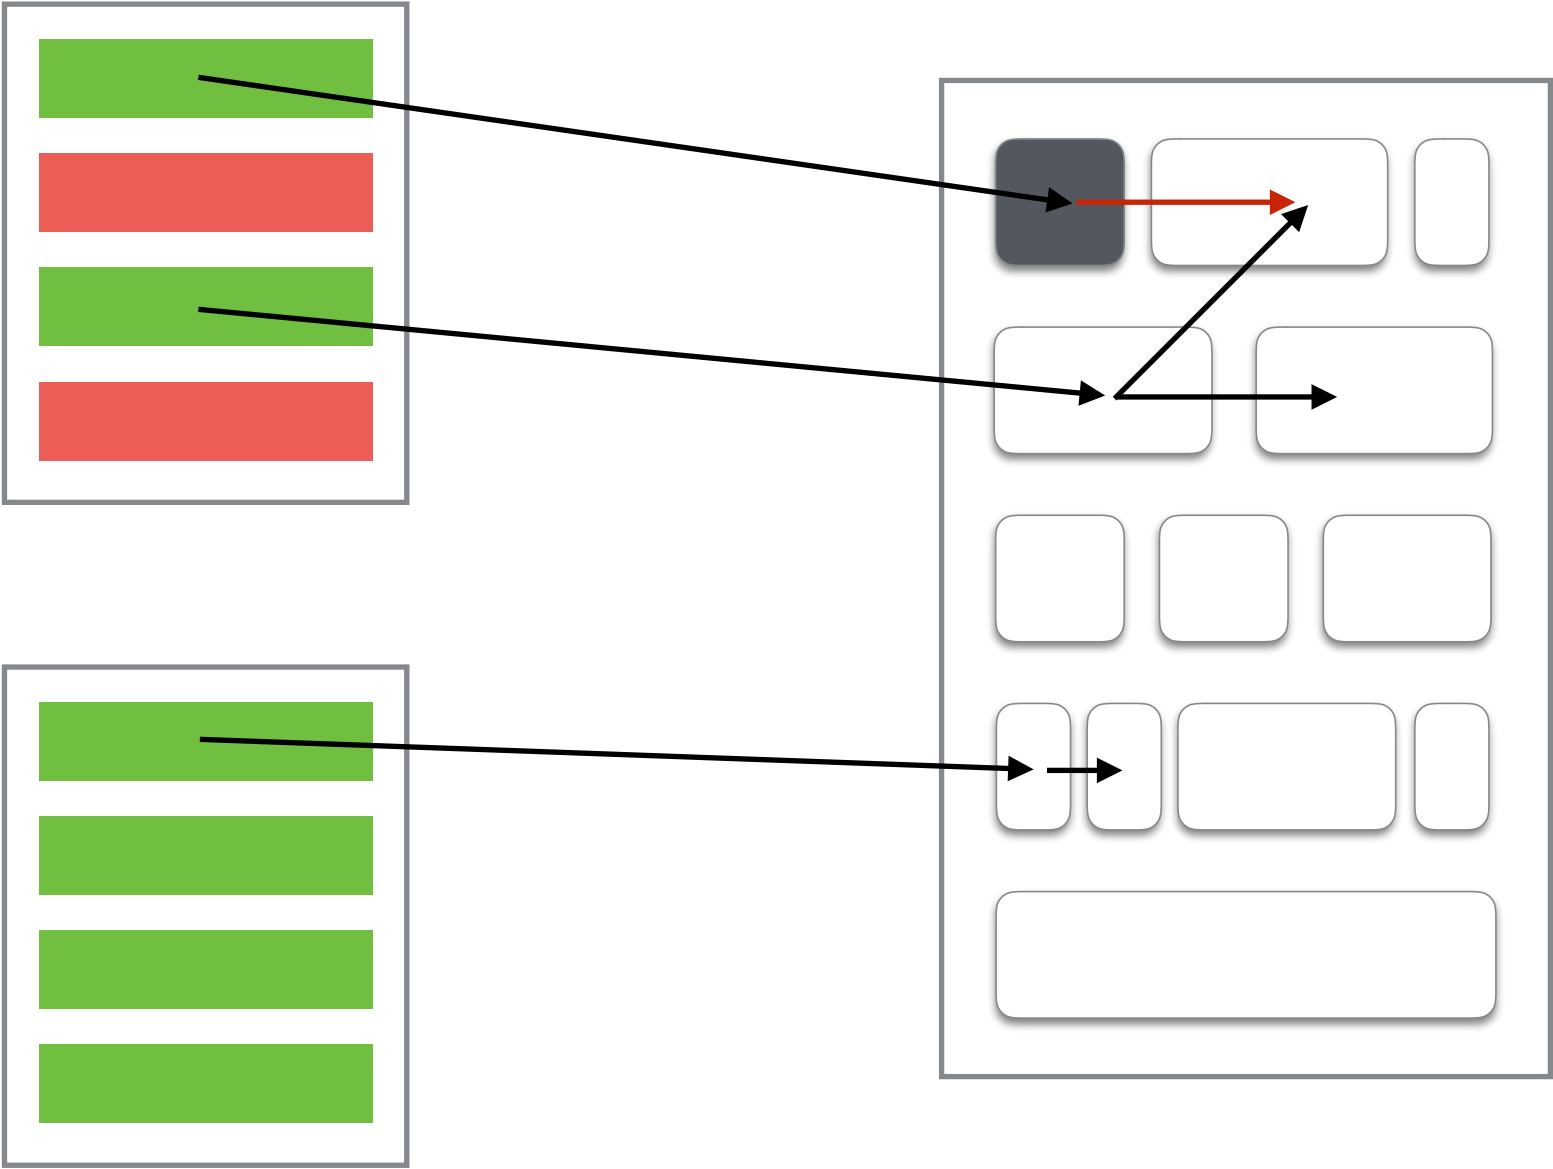


Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase






Modules

MS

Stack

Marking Phase




Modules

MS

Stack

Marking Phase






Modules

MS

Stack


Marking Phase



Modules


MS

Stack





Modules

Sweeping Phase




MS

Stack





Modules

Sweeping Phase




MS

Stack




Modules

Sweeping Phase




MS

Stack





Modules

Sweeping Phase




MS

Stack




Modules

Sweeping Phase




MS

Stack




Modules

Sweeping Phase




MS

Stack





Modules

Sweeping Phase




MS

Stack




Modules

Sweeping Phase




MS

Stack




Modules

Sweeping Phase




MS

Stack





Modules

Sweeping Phase




MS

Stack




Modules

Sweeping Phase




MS

Stack





Modules

Sweeping Phase




MS

Stack




Modules

Sweeping Phase




MS

Stack



Modules

Sweeping Phase



MS Performance




None



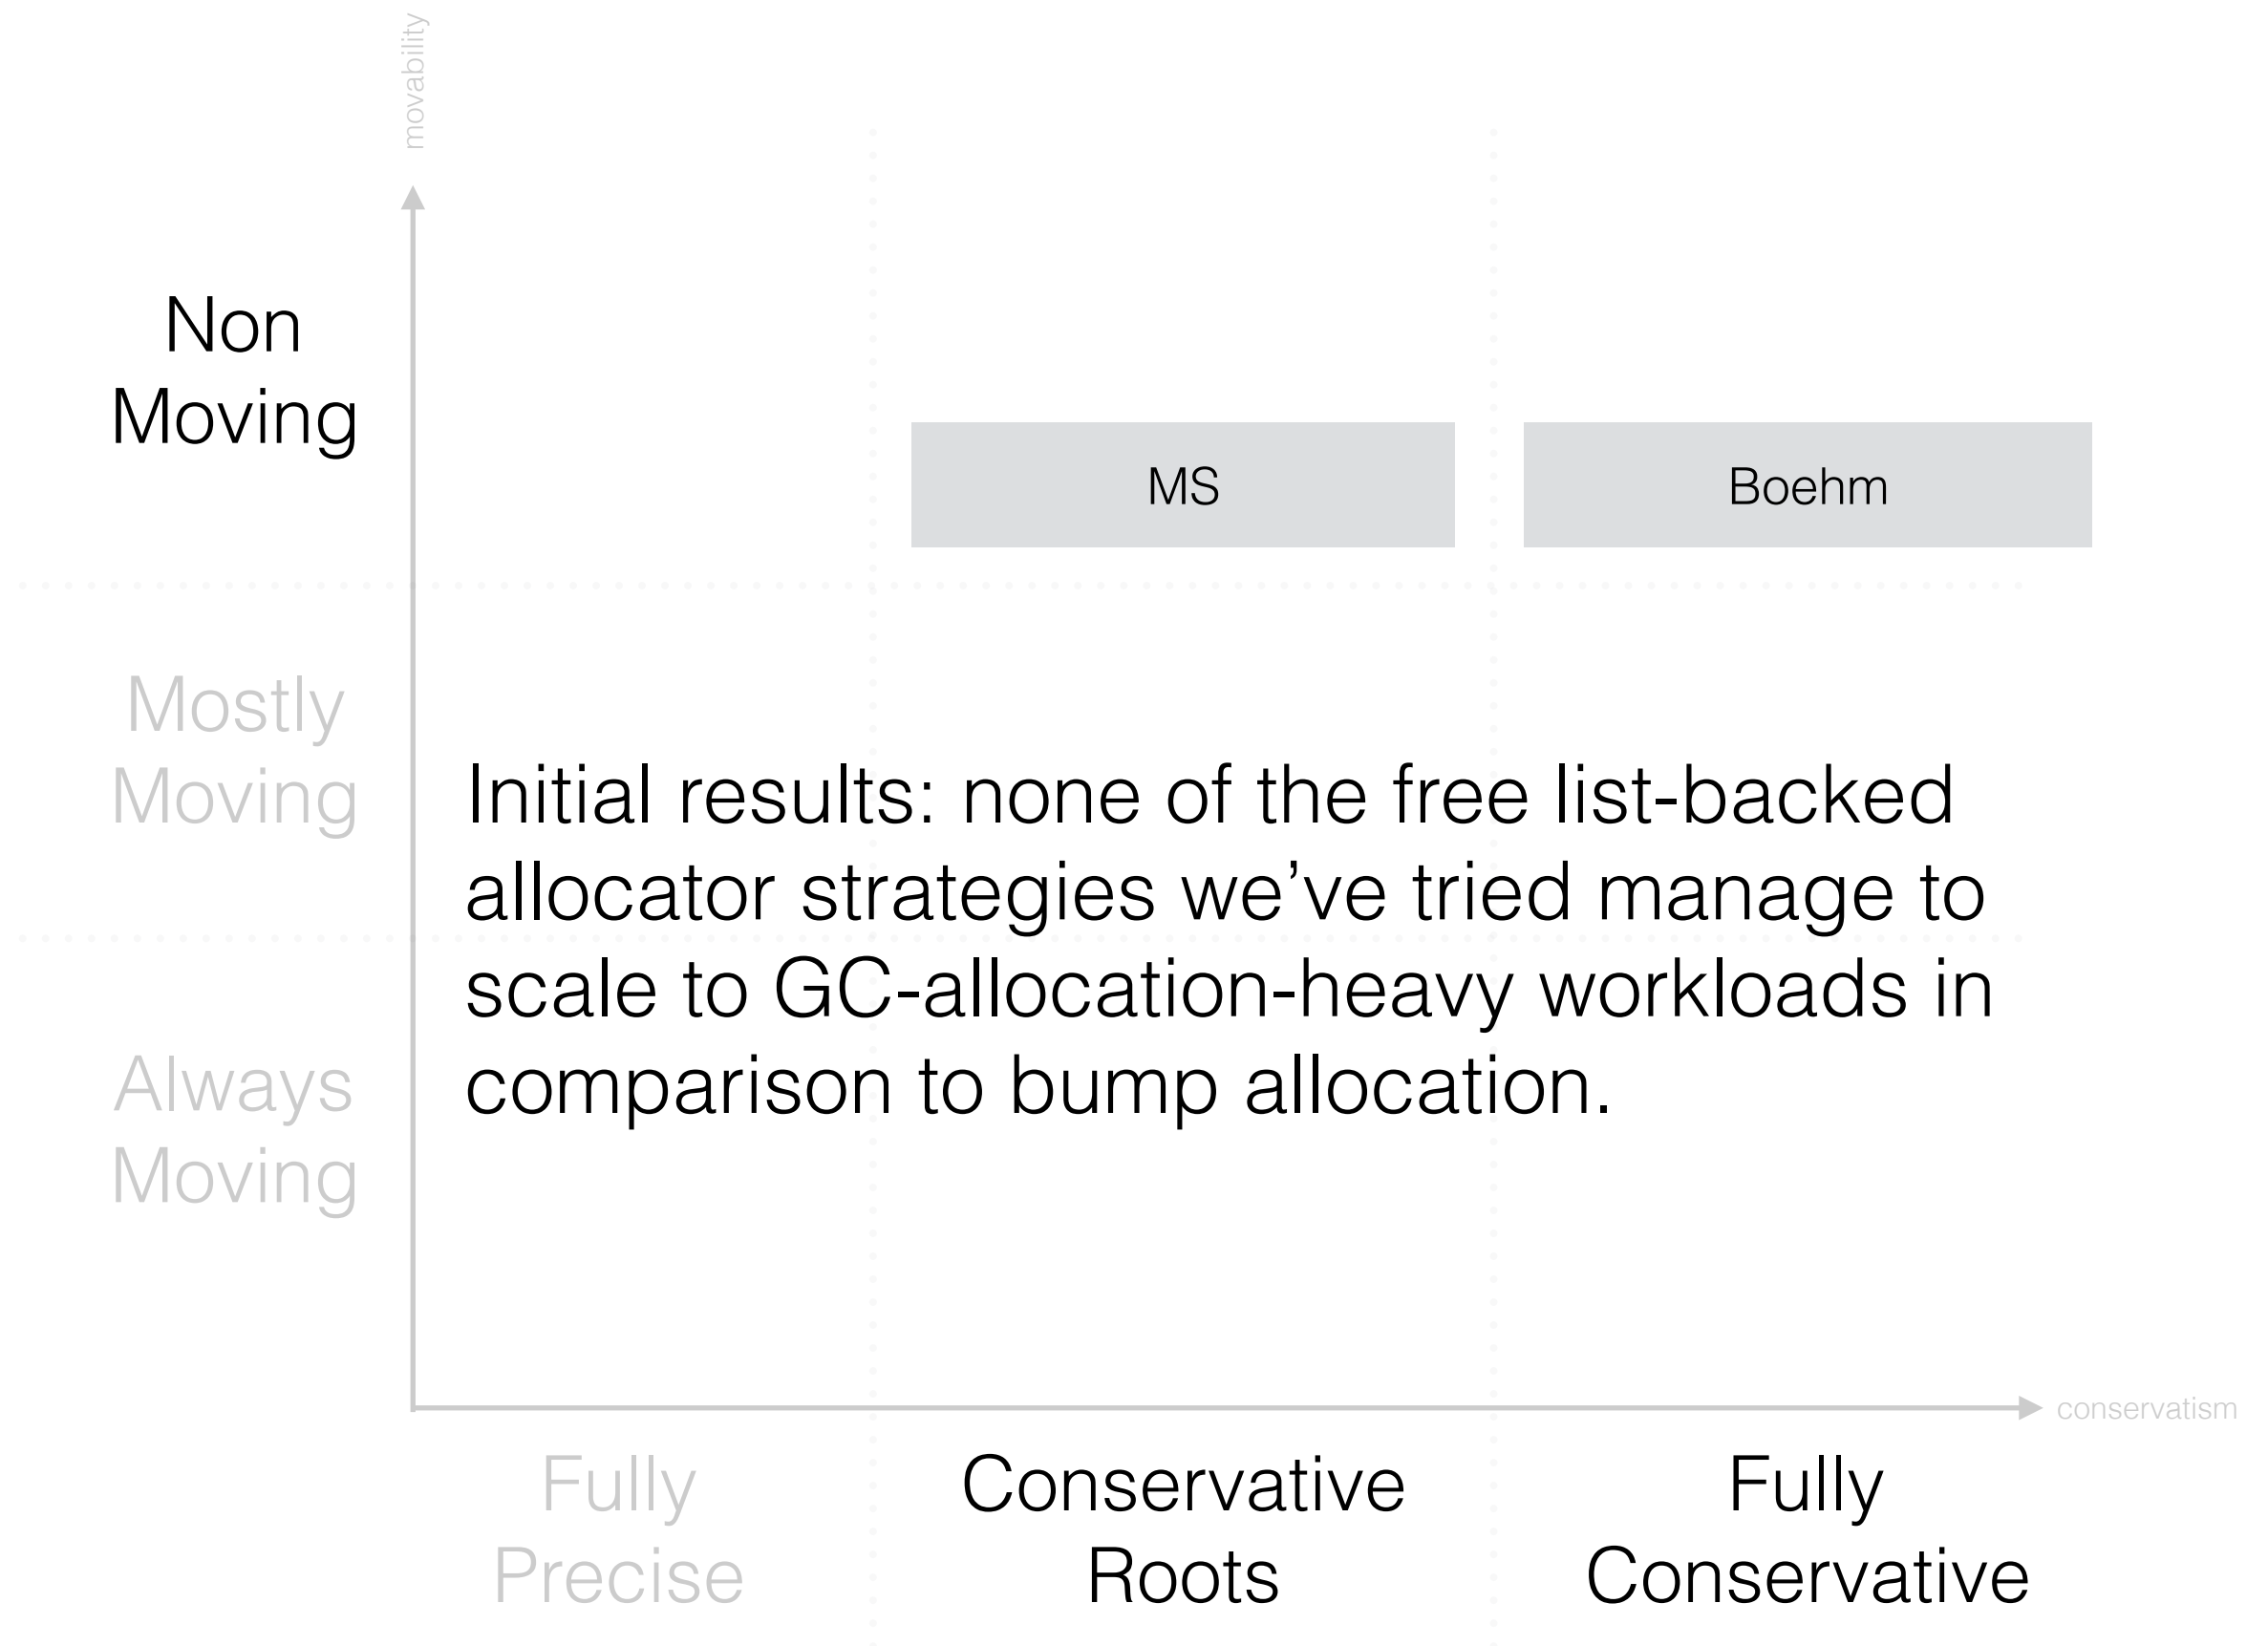
Boehm




MS



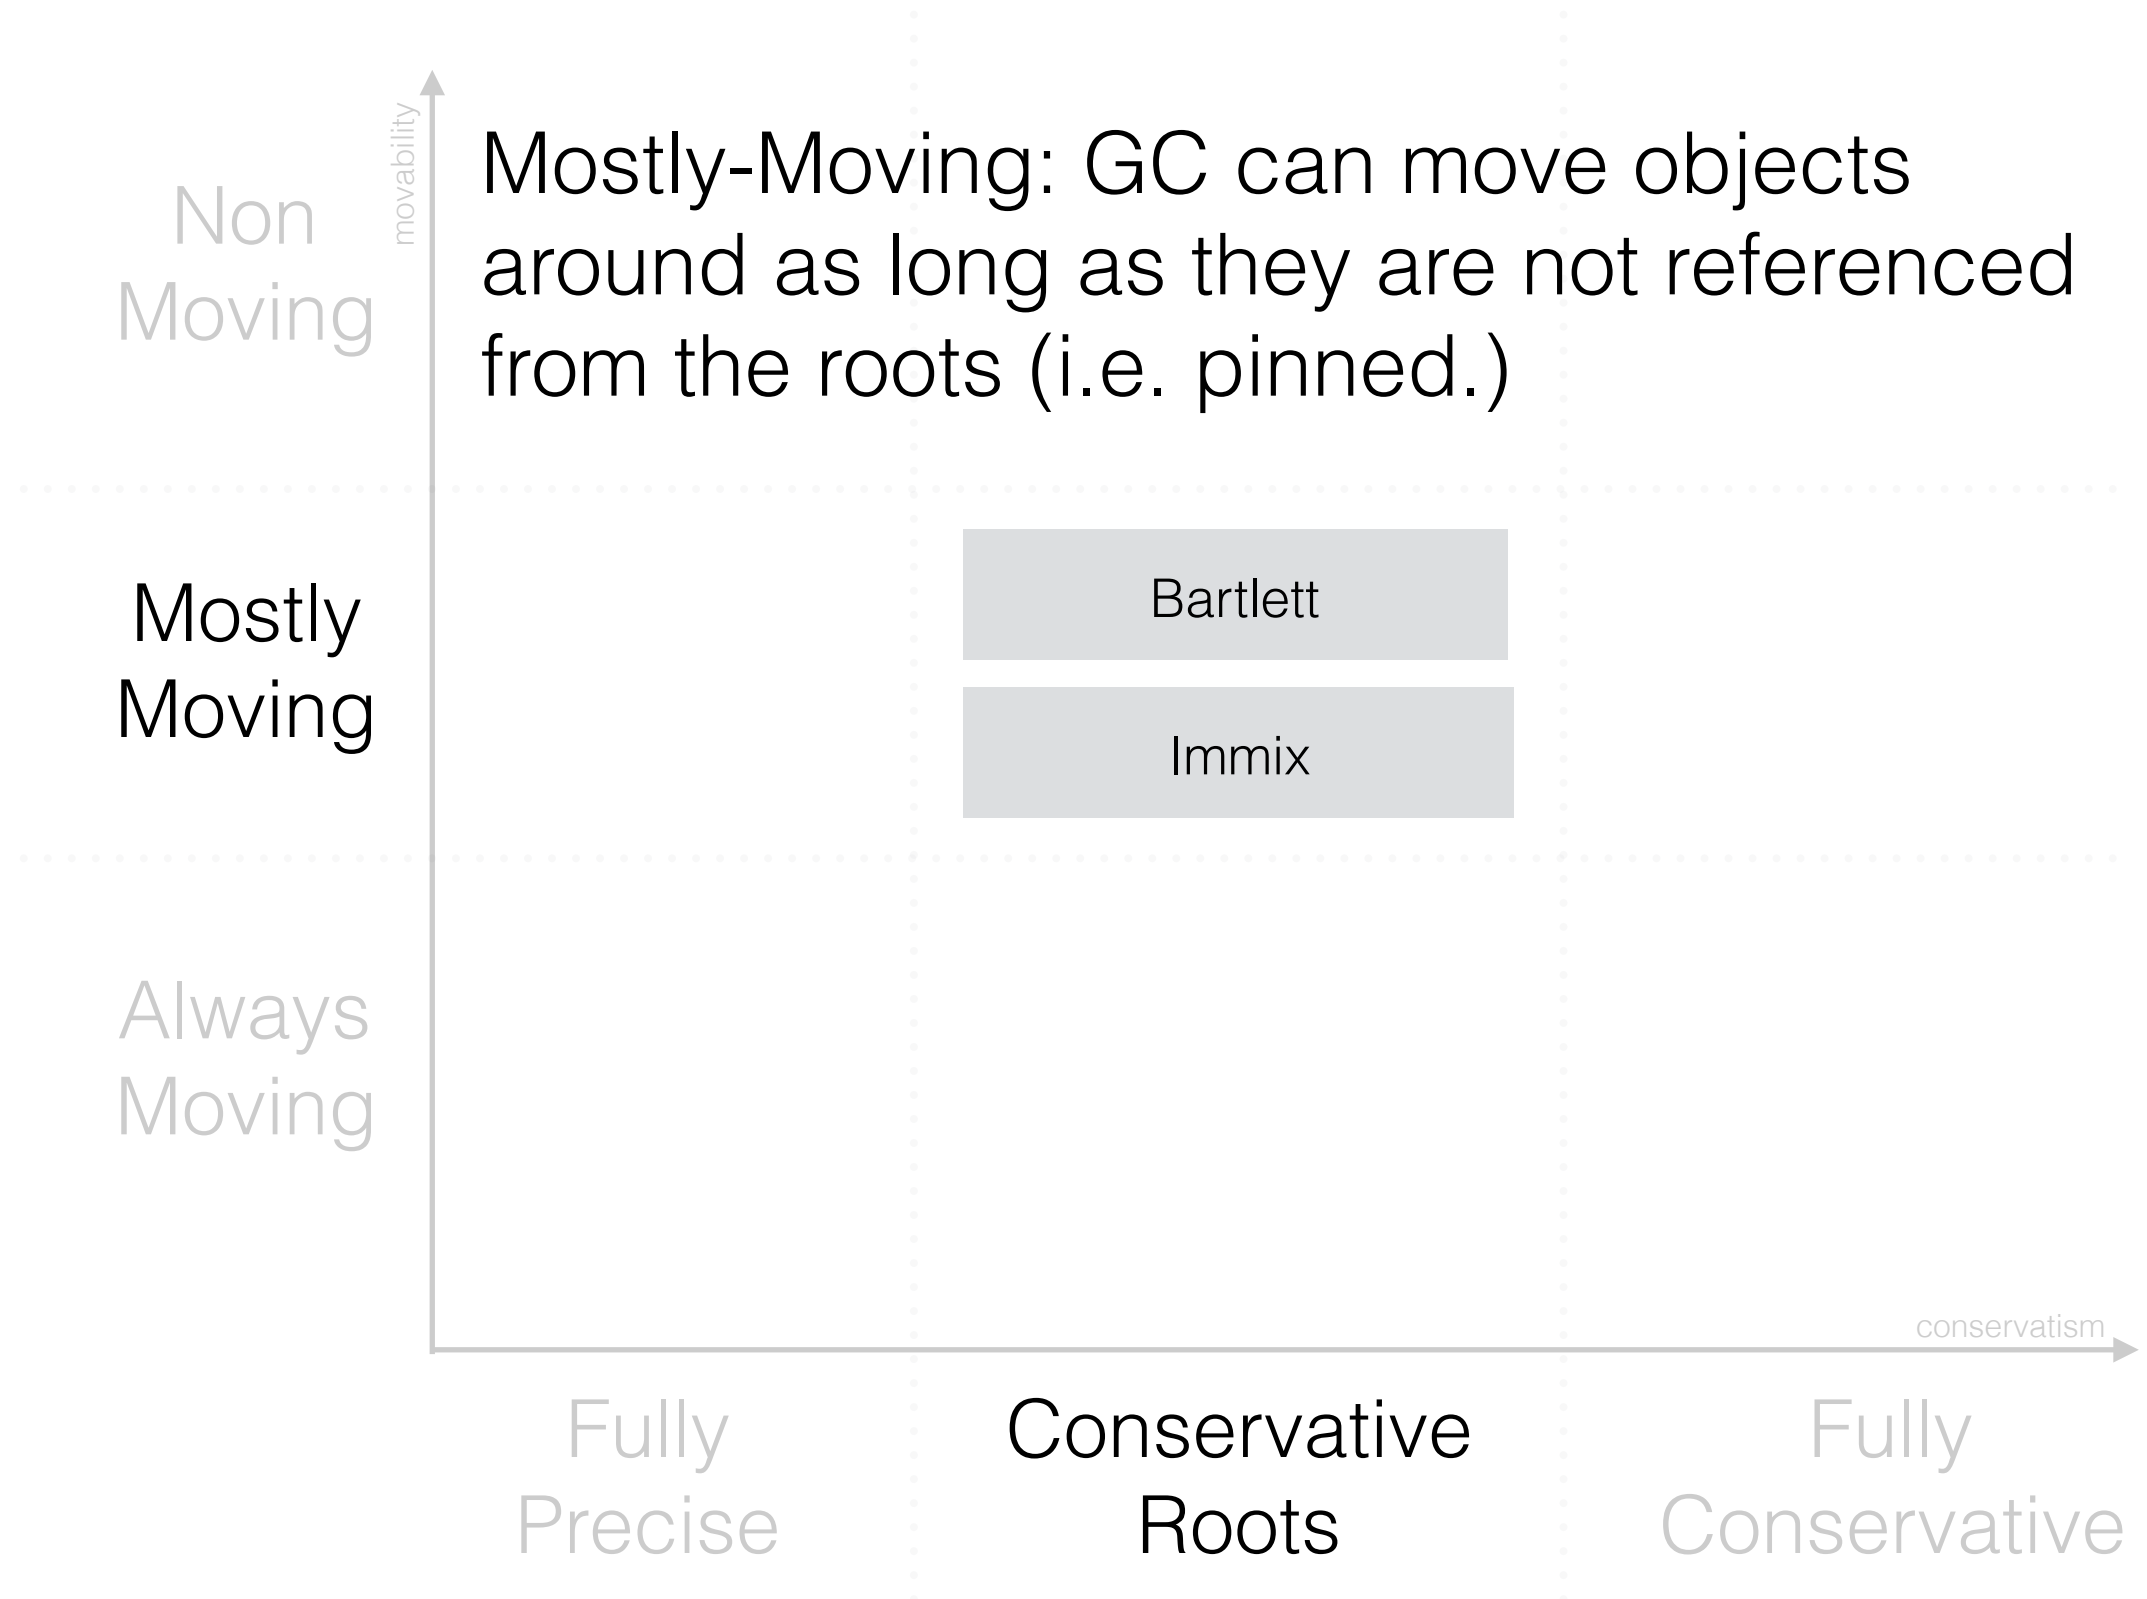
Initial experiments



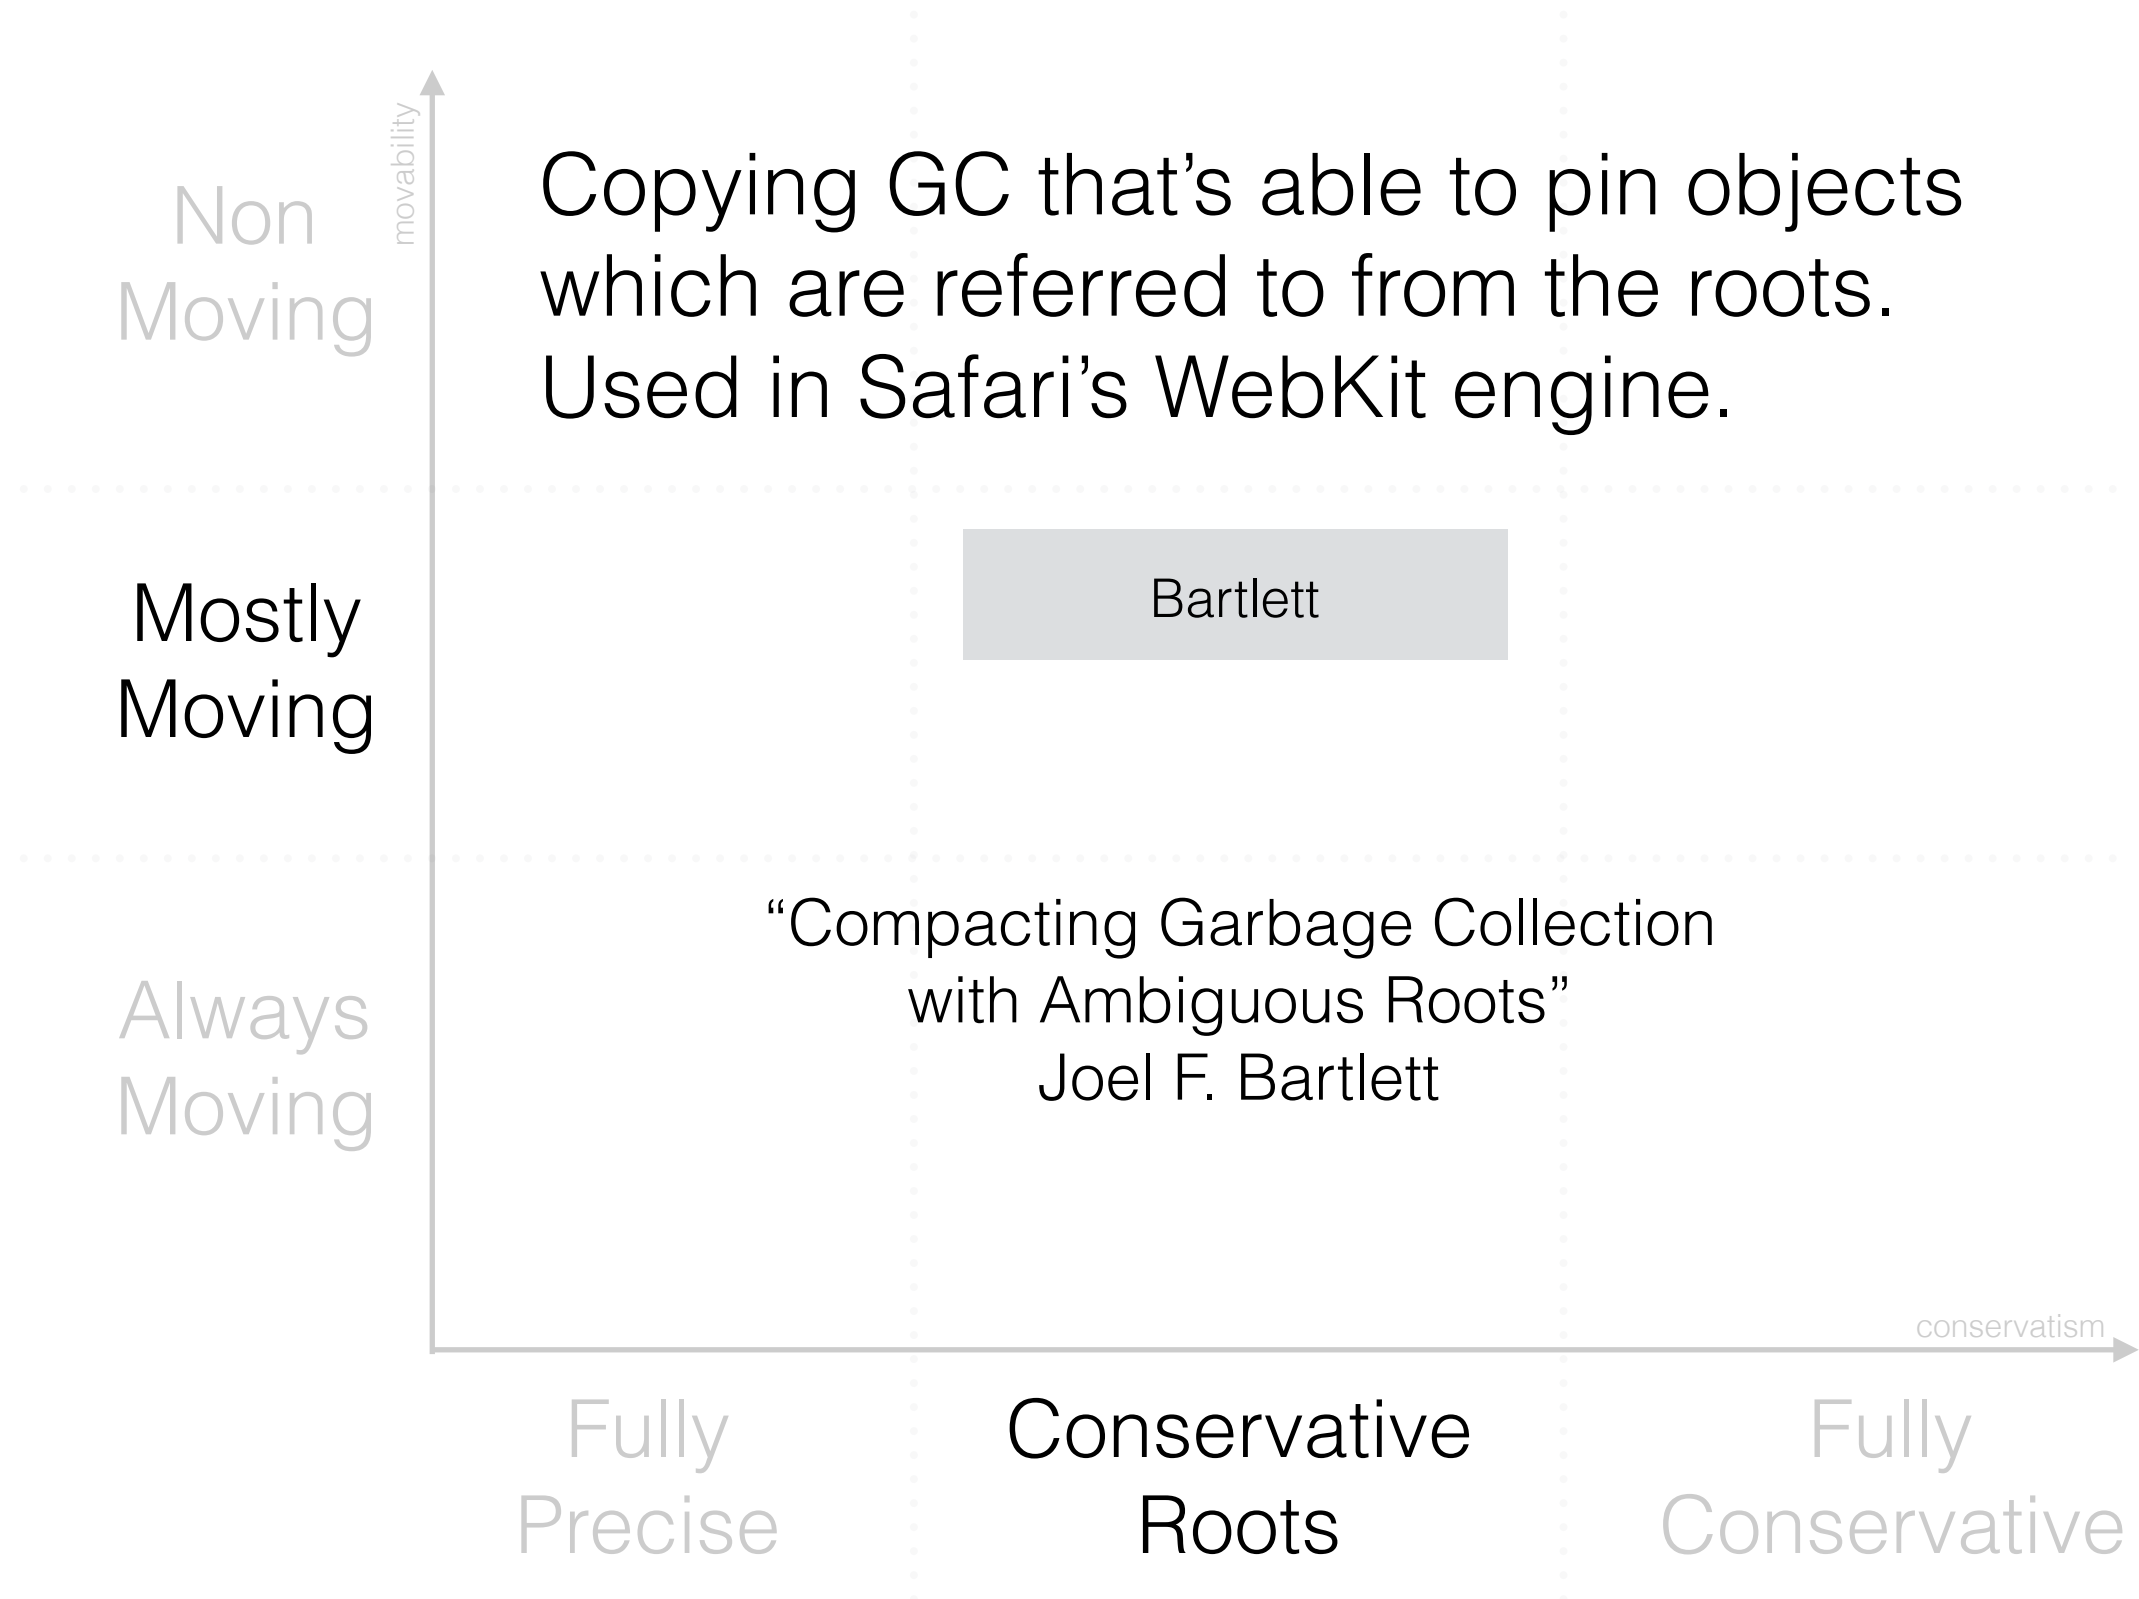
Initial experiments



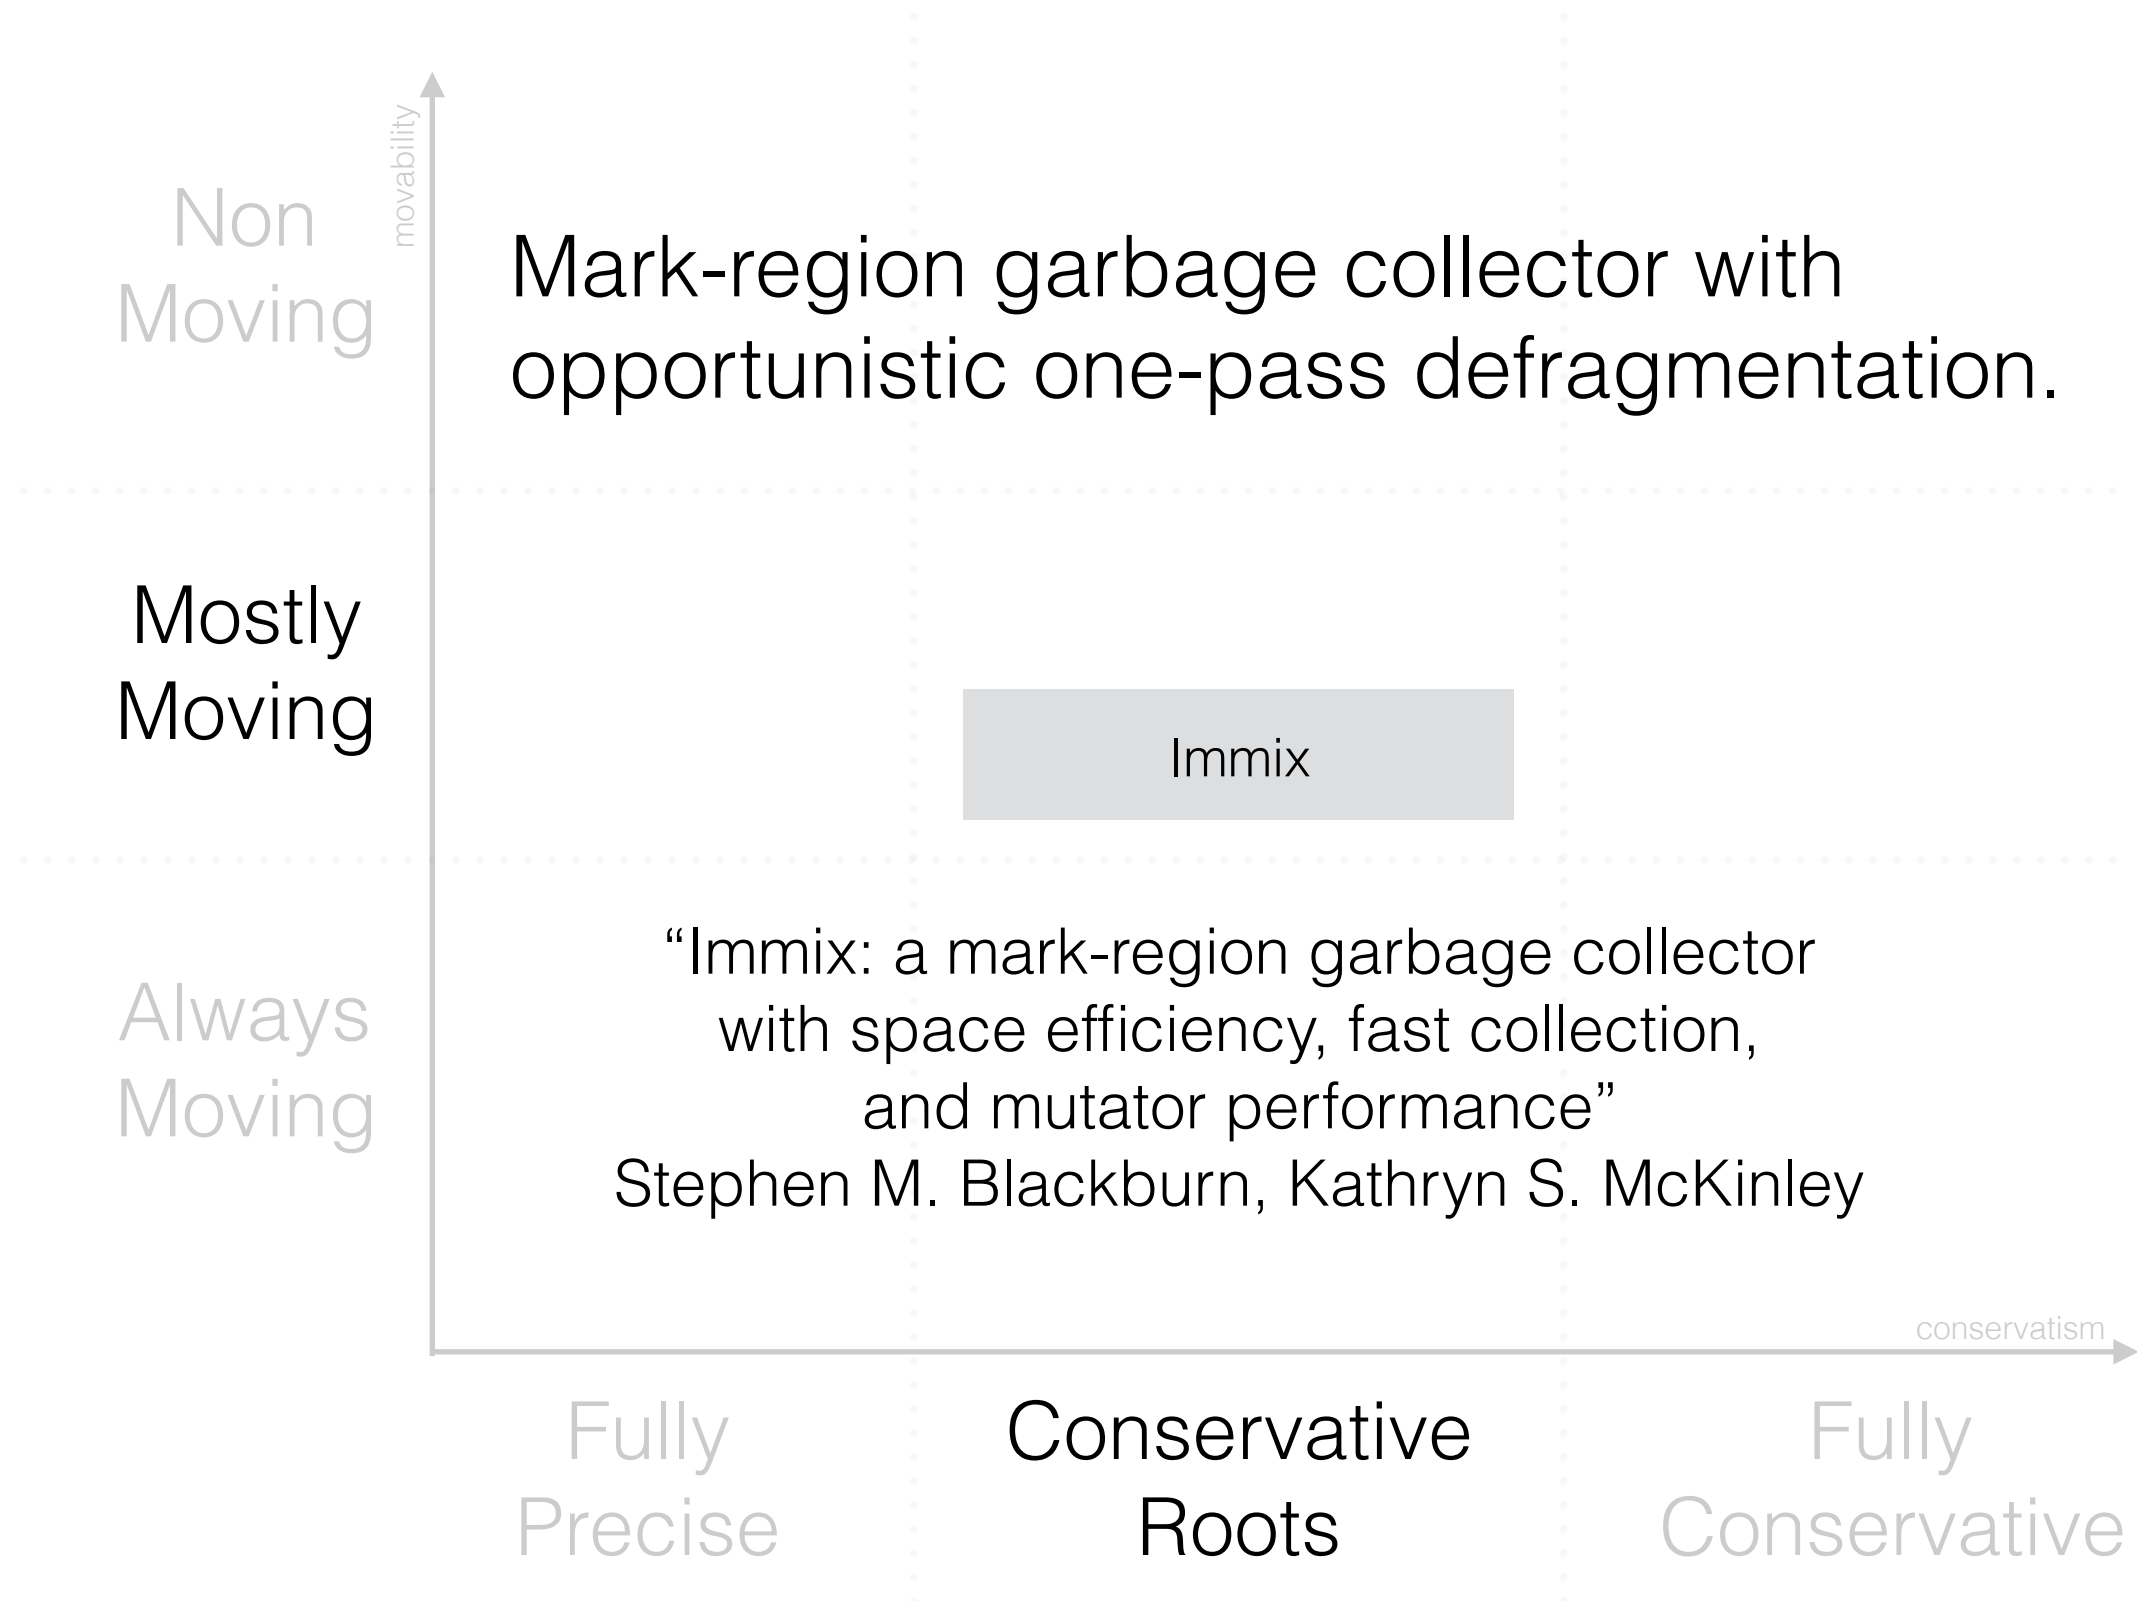
Mostly-Moving Sweet Spot



Bartlett




Immix




Immix

Heap




Immix

Block




Immix


Line



Immix




Immix




Allocation

Immix




Allocation

Immix




Allocation

Immix




Allocation

Immix




Allocation

Immix




Allocation

Immix



Allocation

Immix




Allocation

Immix

Stack

Marking Phase




Modules

Immix

Stack


Marking Phase




Modules

Immix


Stack



Marking Phase




Modules



Immix

Stack



Marking Phase




Modules

Immix

Stack




Marking Phase




Immix


Stack



Modules




Marking Phase




Immix


Stack



Modules




Marking Phase





Immix

Stack





Modules

Sweeping Phase





Immix

Stack




Modules

Sweeping Phase





Immix

Stack




Modules

Sweeping Phase





Immix

Stack




Modules

Sweeping Phase





Immix

Stack




Modules

Sweeping Phase





Immix

Stack



Modules



Sweeping Phase





Immix

Sweeping Phase

Stack





Modules





Immix

Sweeping Phase

Stack




Modules





Immix

Stack



Modules


Sweeping Phase



Immix Performance

Immixon Performance

None Boehm Immix



Immix

- First prototype is coming in Scala Native 0.3
- Opt-in via `nativeGC := "immix"`
- We're going to support Boehm until the immix implementation matures and becomes the default

Bonus features in 0.3

- Sbt testing framework integration
- Initial support for file i/o from `java.nio.*`
- Initial support for zip/jar from `java.util.*`
- Smaller binaries

Questions?